



O3A Core Plugins User Guide

v2.1.6

Copyright 2018 Blue Ripple Sound Limited

Table of Contents

1 Introduction.....	1
1.1 Compatibility.....	1
2 Getting Started.....	2
2.1 Working with O3A.....	2
2.2 Getting Started with Reaper.....	3
2.3 Getting Started with Premiere Pro.....	8
2.4 Getting Started with Pro Tools Ultimate.....	10
2.5 Getting Started with Pyramix.....	12
2.6 Audio for Virtual Reality 360 Video.....	14
2.7 Spotting A Common Issue.....	16
3 O3A Core Plugins.....	17
3.1 O3A Panner.....	18
3.2 O3A Panner - Classic.....	20
3.3 O3A Panner - Hemisphere.....	22
3.4 O3A Panner - Large.....	24
3.5 O3A Panner - Two Channel.....	26
3.6 O3A Panner - Eight Channel.....	29
3.7 O3A Visualiser.....	32
3.8 O3A Visualiser - Hemisphere.....	34
3.9 O3A Colourizer.....	37
3.10 O3A Flare.....	39
3.11 O3A Look.....	41
3.12 O3A Rotation.....	43
3.13 O3A Decoder - Stereo.....	45
3.14 O3A Decoder - Mono.....	47
3.15 O3A Decoder - 5.1 Basic.....	48
3.16 O3A Decoder - O1A, O2A.....	49
3.17 O3A Injector - O1A, O2A.....	50
3.18 O3A Decoder - FuMa.....	51
3.19 O3A Injector - FuMa.....	52
3.20 O3A Beamer.....	53
3.21 O3A Virtual Microphone.....	55
3.22 O3A Meter.....	57
3.23 O3A Meter - Signal.....	59
3.24 O3A Gain.....	60
3.25 O3A Spatial Delay.....	61
4 Technical Notes.....	64
4.1 What is an O3A Stream?.....	64
4.2 What processing can I apply to an O3A stream?.....	64
4.3 Encoding.....	64
4.4 Changes in Version 2.0.....	65
4.5 How does O3A SN3D relate to FuMa and Classic Ambisonics?.....	65
4.6 Exporting First and Second Order.....	65
4.7 Why Third Order?.....	65

1 Introduction

The O3A Core is a set of plugins designed for work with [third order ambisonic \(O3A\) streams](#). These streams can be used to create rich 3D mixes that can be played back in a wide variety of sound systems, including Virtual Reality systems with head tracking.

There are a number of collections of O3A plugins. This one provides a set of essential tools including panning, rotation, visualisation and basic decoding. Other plugin sets available from [Blue Ripple Sound](#) are:

- O3A Decoding
- O3A Decoding - Rapture3D Advanced
- O3A Harpex Upsampler
- O3A Manipulators
- O3A Reverb
- O3A Upmixers
- O3A View

1.1 Compatibility

O3A streams require audio busses with at least 16 channels, which (at the time of writing) most Digital Audio Workstations (DAWs) cannot handle. Some are limited at 2 channel stereo and many are limited at the 8 channels used for 7.1 surround mixes. **The O3A plugins *will not work correctly* with these DAWs**, which may even crash.

1.1.1 AAX

Most of the O3A plugins exist as AAX versions which are compatible with [Pro Tools Ultimate](#) or Pro Tools HD 12.8.2 and later and plugin support (or not) is indicated in this text. Pro Tools keeps careful track of exactly what stem formats are present on tracks and busses; this is a powerful feature which helps avoid mistakes and ensure plugins are used in the right places. However, where formats are not available in Pro Tools, corresponding plugins are generally not available. To avoid this in some cases, some plugins "misuse" stem formats with matching channel counts. Cases like this are described for individual plugins below in the text.

1.1.2 VST2

The plugin library works as a "shell" plugin. This means that a number of individual audio plugin effects are provided by a single library file. Some VST hosts may have a slightly different way of managing these plugins to ordinary ones. For instance, in [Max/MSP](#) the vst~ plugin uses "subname" messages to specify the individual plugin within the library.

Most VST2-compatible DAWs (such as [Reaper](#)) have a plugin "path", which is a list of directories which will be searched for VST2 plugins. You may need to change this path to point at the location of the plugins, or move the plugins there. By default, these plugins are installed into /Library/Audio/Plug-Ins/VST on macOS. Various directories may be used on Windows, but C:\Program Files\Steinberg\VST2 is not uncommon.

1.1.3 Buffering

Some of the O3A plugins use internal buffering with a length of 128 samples. For smooth CPU load, you may wish to ensure your DAW buffer size is a multiple of this.

2 Getting Started

2.1 Working with O3A

When working with stereo, you normally "pan" a mono sound to stereo and "mix" it onto a stereo bus. With O3A, it is similar but there is an extra bus and "decoding" step:

1. Panning - mono sources are panned to 16 channel O3A. The [O3A Panner](#) plugin can be used for this.
2. Mixing - multiple streams of 16 channel O3A can be mixed together using your DAW's mixing engine. No special plugin is needed here, as long as the DAW can handle 16 channel busses. You end up with a final mix, as an O3A stream.
3. Decoding - during actual playback (and you probably want to listen while you work!) you "decode" the O3A stream for whatever speaker layout you are actually listening with. One of the huge advantages of ambisonic techniques is that the final O3A mix above is independent of its final playback environment, so you can switch between stereo speakers, headphones, 5.1, 22.2 or whatever. One of the simplest of these decoders is provided by the basic [O3A Decoder - Stereo](#) plugin.

These are the basics. You can also manipulate O3A streams, for instance to [rotate](#) the soundfield. Plugins that do this have 16 channel O3A as both an input and output. There are some [technical notes on O3A streams](#) at the end of this document.

If you are also working with first order B-Format using the SN3D/ACN convention you can mix this in too, because the first four channels of O3A correspond to the four channels of first order. To reduce an O3A stream to first order, you can simply take these first four of the sixteen O3A channels and throw away the rest.

2.2 Getting Started with Reaper

This section gives an introduction to use of the VST2 version of the O3A plugins with [Reaper](#).

We are assuming that you know how to use Reaper already. A number of simple Reaper example projects are included in standard installations, along with simple explanatory text files.

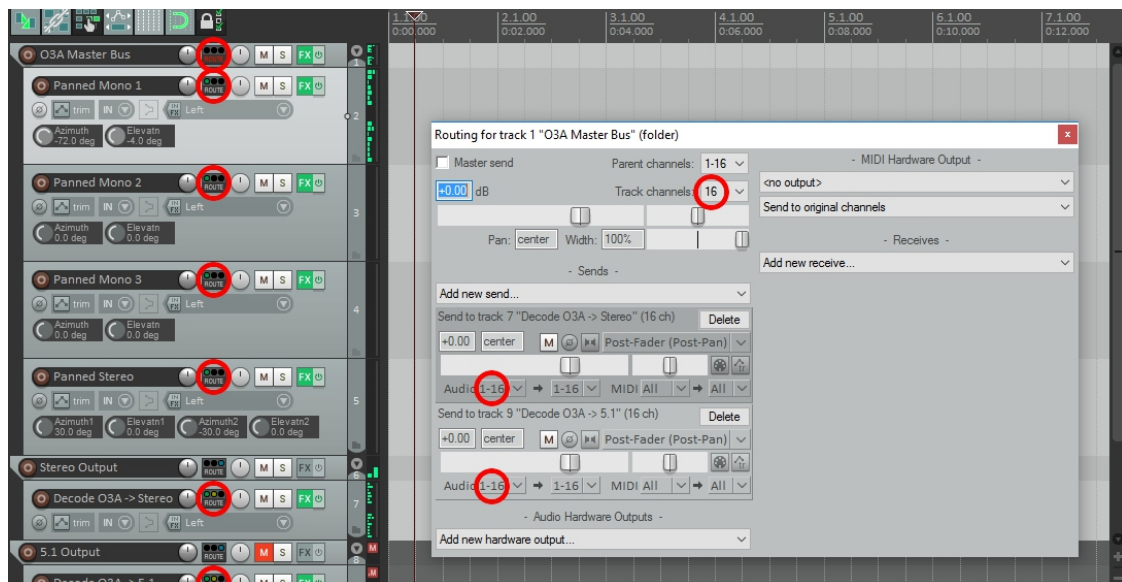
2.2.1 Shell Plugins

The O3A VST plugin libraries are VST2 "shell" plugins which each provide a number of plugin effects. If you find your list of O3A effects is short, with names like "O3ACore", this means that Reaper has not scanned the content of the libraries.

To fix this, go into the Reaper "Options" menu and select "Preferences..." to bring up the preferences screen. Then, under "Plug-ins", find "VST". On the VST tab, please ensure that "Get VST names/types when scanning" is checked and then click "Clear cache/re-scan".

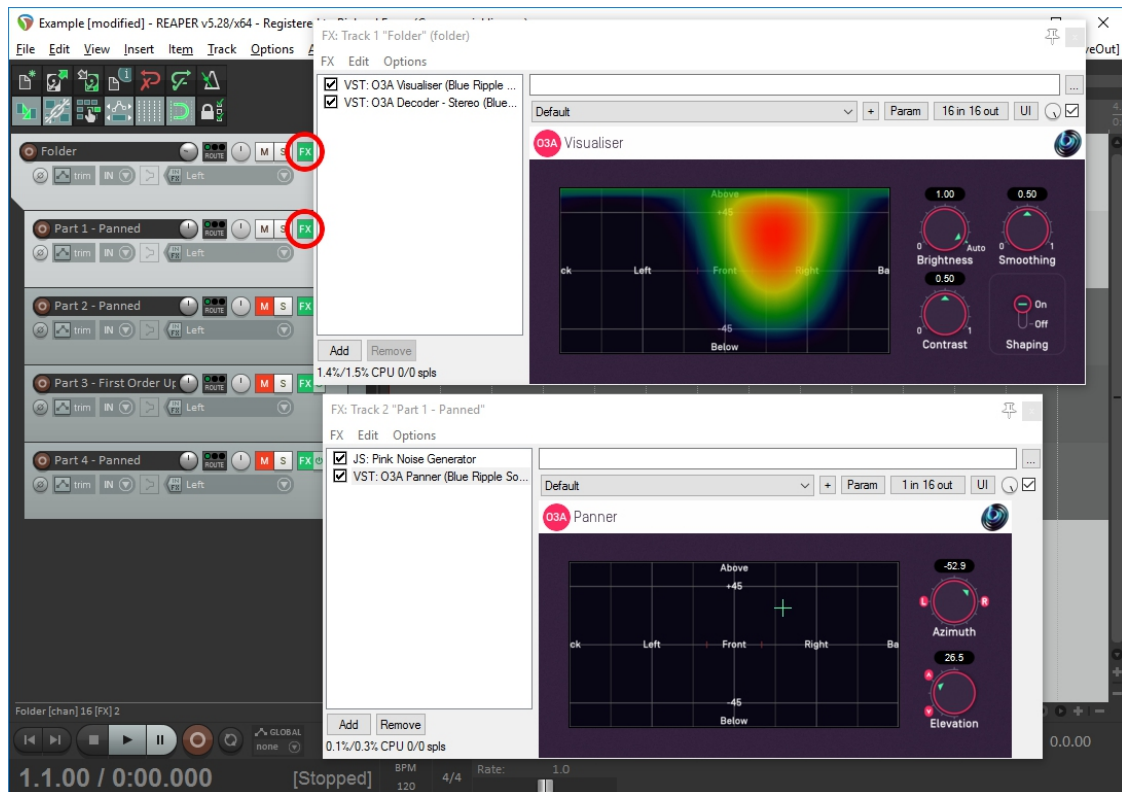
2.2.2 Basics

In Reaper, whenever you create a track or send to be used with [O3A streams](#), make sure it has at least 16 channels. A few plugins need even more (e.g. a decoder for Hamasaki 22.2).



2.2.3 Layout

A practical way to set up projects in Reaper is to create one "folder" track which takes O3A and decodes it, and a number of individual tracks within that folder which feed O3A to it. A slightly more elaborate way to do this, supporting a choice of decoders, is described in the Reaper example projects.



On the folder track, add the following plugins *in this order* to the FX chain:

- [O3A Visualiser](#)
- [O3A Decoder - Stereo](#) (or whichever decoder best suits your speaker rig).

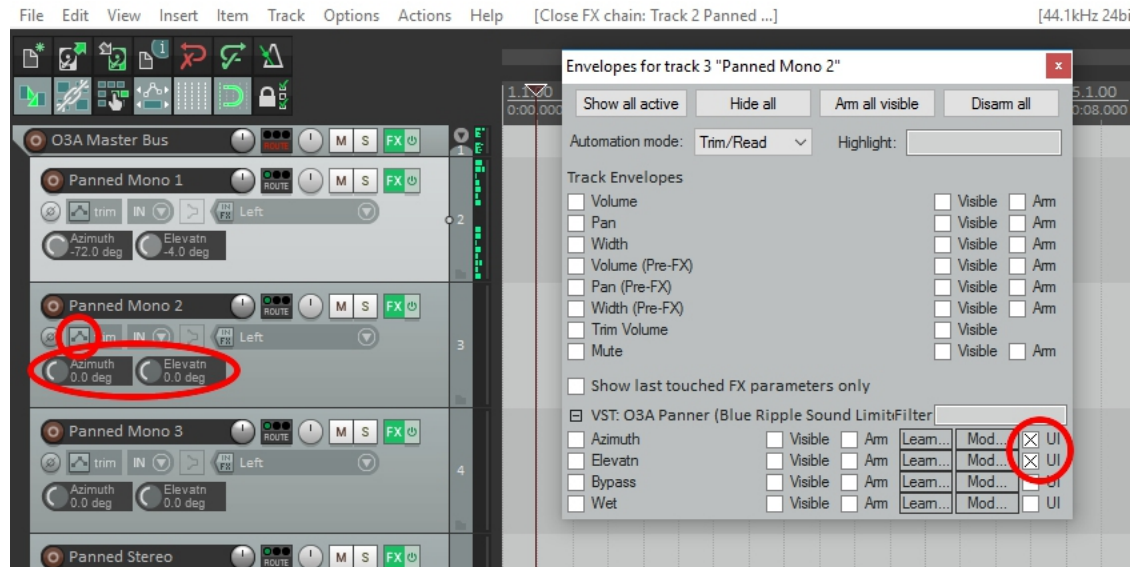
This means that after all the audio from the individual parts has been mixed together it will be visualised in 3D and then converted to stereo for monitoring.

On tracks for mono parts that are to be panned, add the following:

- An [O3A Panner](#) plugin.

Do *NOT* use the built-in Reaper stereo panner with O3A!

It is also possible to add azimuth and elevation controls to Reaper on these channels. To do this, add the panner as above and then use the "Trim" button and cross the "UI" boxes for the azimuth and elevation controls.



Parts that are already first order ambisonic B-Format can typically be injected directly without issue, although processing can help (e.g. the O3A First Order Injector plugin from the O3A Upmixers set, or the O3A Harpex Upsampler). Of course, other O3A material can be injected directly. FuMa material needs to be converted using the [O3A Injector - FuMa](#) plugin.

2.2.4 Processing O3A streams

Other plugins (for instance the [rotation](#) plugin) manipulate the O3A stream. The stream needs to be in O3A format (for instance by panning) before such processing can be used.

Applying processing to the O3A stream can be confusing at first. We suggest you spend a few minutes getting familiar with the [O3A Visualiser](#) plugin and put visualisers before and after your processing. You can then see what is going on, which can be really helpful!



Some plugins have O3A as an input but output something else. In these cases, make sure you only chain together processors in a combination that make sense. For instance, the [O3A Beamer](#) plugin produces mono which you might want to run through a mono distortion plugin before panning it back into O3A using the [O3A Panner](#).

2.2.5 Sends

Effect sends can be used in the normal way. Just be careful about which busses are expecting mono or O3A as input and output and be sure to check your channel counts.

Again, do not use the Reaper panner for the send (leave it set to the centre) or the spatial image will be damaged. For similar reasons, make sure the project's "Pan Law" (in Project Settings) is set to 0dB, or avoid use of Post-Fade sends.

2.2.6 Mastering

2.2.6.1 Keep Your O3A Master Mix

Keep your Reaper project files and inputs if you like, keep your favourite stereo and 5.1 downmixes etc. But whatever you do, *keep a copy of the 16 channel O3A mix*. This may not make a difference to you now, but when someone asks you for a 22.2 (or other) downmix in a few years time, you will not need to fake height information from the 5.1 mix, or go searching for a missing file needed by the original Reaper project. Instead, you will be able to produce the downmix from the ambisonic B-Format (i.e. O3A) master using whatever tools are in vogue then.

In Reaper, probably the easiest way to grab the O3A mix is to switch off your decoder (in your folder track if you have followed the advice above) and then use the "Render..." menu item. You will probably need to set the Channels to "16" by typing the number in and use the "Stems (selected tracks)" option rather than what Reaper considers the master mix, in which case you will also need to have selected your folder track. If you want to check the output, you can always load it into a new Reaper session and run it through your favourite decoder.

2.2.6.2 Other Mixes

If you want to produce a number of different format outputs from your O3A master (stereo, headphones, 5.1 etc.) you could use a number of projects that use the O3A master file.

Another option is to place a number of decoders on your decoding/folder track and only switch on the one you want for the particular rendering.

However, our preferred option is to send from the O3A master bus to one of a number of separate decoding tracks which can be freely muted or unmuted. This is the approach taken in the example Reaper projects provided.

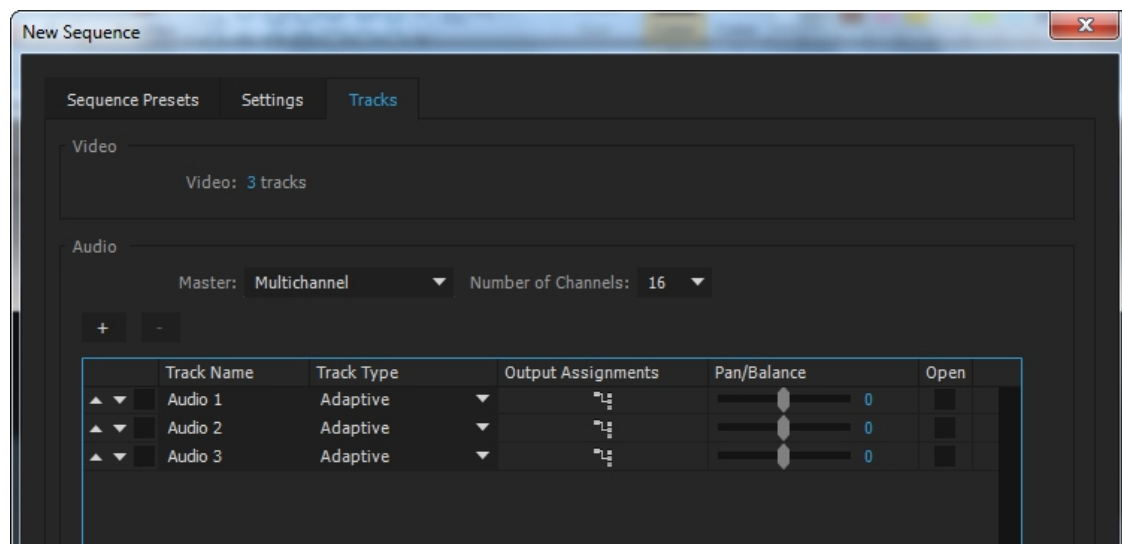
2.3 Getting Started with Premiere Pro

In Adobe Premiere Pro, as with Reaper, it is essential to make sure that your audio tracks have 16 channels (or more if you are using plugins that need them).

2.3.1 Sequence Settings

In particular, when creating a New Sequence, go to the "Tracks" tab and make the following changes:

1. The "Master" Track should be set to "Multichannel".
2. The "Number of Channels" should be set to "16", or as many as you need.
3. Each audio "Track Type" that is going to carry O3A audio (or suchlike) should be set to "Adaptive" or "Adaptive Submix".



Premiere Pro includes some VR Sequence Presets that support ambisonics. However, these only set the number of channels to 4 by default, which is correct for first order ambisonics, but not for third order. 16 channels are needed by these plugins, so please check the "Number of Channels" as above, or plugins may not work and spatialisation quality may be seriously compromised.

2.3.2 Importing Audio

When importing audio with many channels, Premiere Pro may place each channel of the audio onto a different track rather than keeping the spatial mix together. If this happens, find the clip, right click and select "Modify" and then "Audio Channels..." and do the following before trying again:

1. Set the "Clip Channel Format" to "Adaptive".
2. Check the "Active Channels Per Clip" matches the audio (for instance, 4 for first order ambisonics or 16 for third order ambisonics).
3. Set the "Number of Audio Clips" to 1.

You can also make this the default behaviour by going into Premiere Pro's "Preferences...", finding the "Audio" section and setting the "Multichannel Mono Mode" item to "Adaptive".

2.3.3 Compatibility

At the time of writing, these plugins are believed to work with the current version of Adobe Premiere Pro (v12.1.0). They do not work with preceding versions.

The "O3A View Sync" plugin from the O3A View library is not fully compatible with the current version of Premiere Pro. Necessary VST2 timing information is currently not available in Premiere, which means that video playback cannot synchronize correctly.

The Premiere list of "Available Plug-Ins" only shows one plugin from each of the plugin libraries, but they all should be available for use.

Plugin presets are not available in Premiere Pro.

2.4 Getting Started with Pro Tools Ultimate

This section gives an introduction to use of the AAX version of the O3A plugins with [Pro Tools Ultimate](#), which was previously known as Pro Tools HD. Ambisonic support has been available in this edition of Pro Tools (but not regular Pro Tools) since version 12.8.2.

These notes assume you are familiar with Pro Tools Ultimate already. A number of simple Pro Tools Ultimate example projects are included in standard installations, along with simple explanatory text files.

2.4.1 Stem Formats

When effects are used on tracks, Pro Tools Ultimate takes account of changes in the stem format in use, which is generally very convenient when working with ambisonics. In particular, effect lists shown will be appropriate to the current stem format you are using.

For instance, if you have a mono track that you wish to pan to O3A, you will find the insert effect list includes the [O3A Panner](#) (under "Sound Field"). Introducing this effect as an insert will change the track stem format from this point to third order ambisonics (O3A) automatically. Mono effects can still be used *before* the panner, but after panning the effect list will include a different set of plugins appropriate for a third order ambisonic stream, such as the [O3A Visualiser](#) and O3A decoders. If you decode to stereo, the stem will change to stereo and so on.

2.4.2 Stem Formats that Pro Tools Ultimate does not support directly

However, Pro Tools Ultimate does not have direct support for some of the more unusual stem formats used by the VST versions of the plugins. Consequently:

- Some of the plugins here do *not support* AAX. For instance, plugins requiring the Hamasaki 22.2 stem format are not included in the AAX version of the plugin libraries.
- Some plugins "misuse" other stem formats with matching channel counts. For instance, the [O3A Decoder - FuMa](#) plugin outputs ambisonics using the FuMa convention, but labels its output stem format as ordinary ambisonics using the SN3D convention, which is not really correct. As a result, the Pro Tools effect lists shown will not be correct. Similarly, the O3A Decoder - IMAX 6.0 plugin outputs audio targetting the IMAX 6.0 layout, but labels its output stem format as standard 6.0, which does *not* have exactly the same channels.


Where you are using standard Pro Tools Ultimate stem formats only, these issues are not relevant. But, if you are using stem formats that Pro Tools does not support directly, care is needed.

Please see the manuals for details of AAX support for each plugin.

2.4.3 Ambisonic Tracks

Pro Tools Ultimate stem formats make it very straightforward to organise an ambisonic project. Often, tracks do not need to be created in special way as simply including an O3A panner or other ambisonic encoder will result in the track stem format changing automatically, e.g. from mono.

However, if you wish to record audio directly in an ambisonic format, this can be handled by creating an ambisonic track. Pro Tools simply provides options here on the track creation screen.



Like these plugins, Pro Tools normally uses the [SN3D/ACN](#) ambisonic format. You can convert from FuMa to SN3D using the [O3A Injector - FuMa](#) plugin.

2.4.4 Organising Projects

A recommended way to organise a project is to route all of your ambisonic material (panned or otherwise) to a single ambisonic "O3A Master" bus, set to third order ambisonics in much the same way as tracks (see above).

This gives a convenient place to bring the whole O3A mix together and send it to a number of tracks for visualisation (e.g. with [O3A Visualiser](#) or [O3A Flare](#)), metering and decoding for specific output formats.

For instance, you might create a "Stereo Decoding" track including the [O3A Decoder - Stereo](#) plugin, so the final stem format is stereo, ready for monitoring. You might also set up a "5.1 Decoding" track in parallel and mute the one you are not using - and so on for other final output formats.

Effect sends can be used in the normal way and Pro Tools will help you match up the stem formats in use.

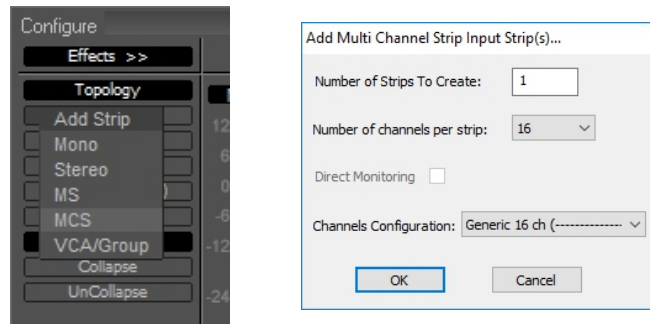


2.5 Getting Started with Pyramix

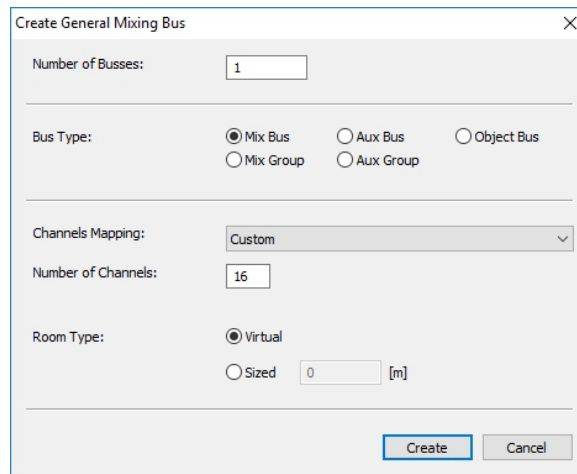
This section gives an introduction to use of the VST2 version of the O3A plugins with [Pyramix](#). We are assuming you know your way around Pyramix already and have configured it to find the VST plugins. Note that modern Pyramix requires the 64bit version of the plugins.

2.5.1 Setup

[O3A streams](#) are most easily handled using Multi Channel Strips (MCS). When creating these, set the channel count to 16 (or even more, e.g. for Hamasaki 22.2).



Custom 16 channel mix busses can also be used.

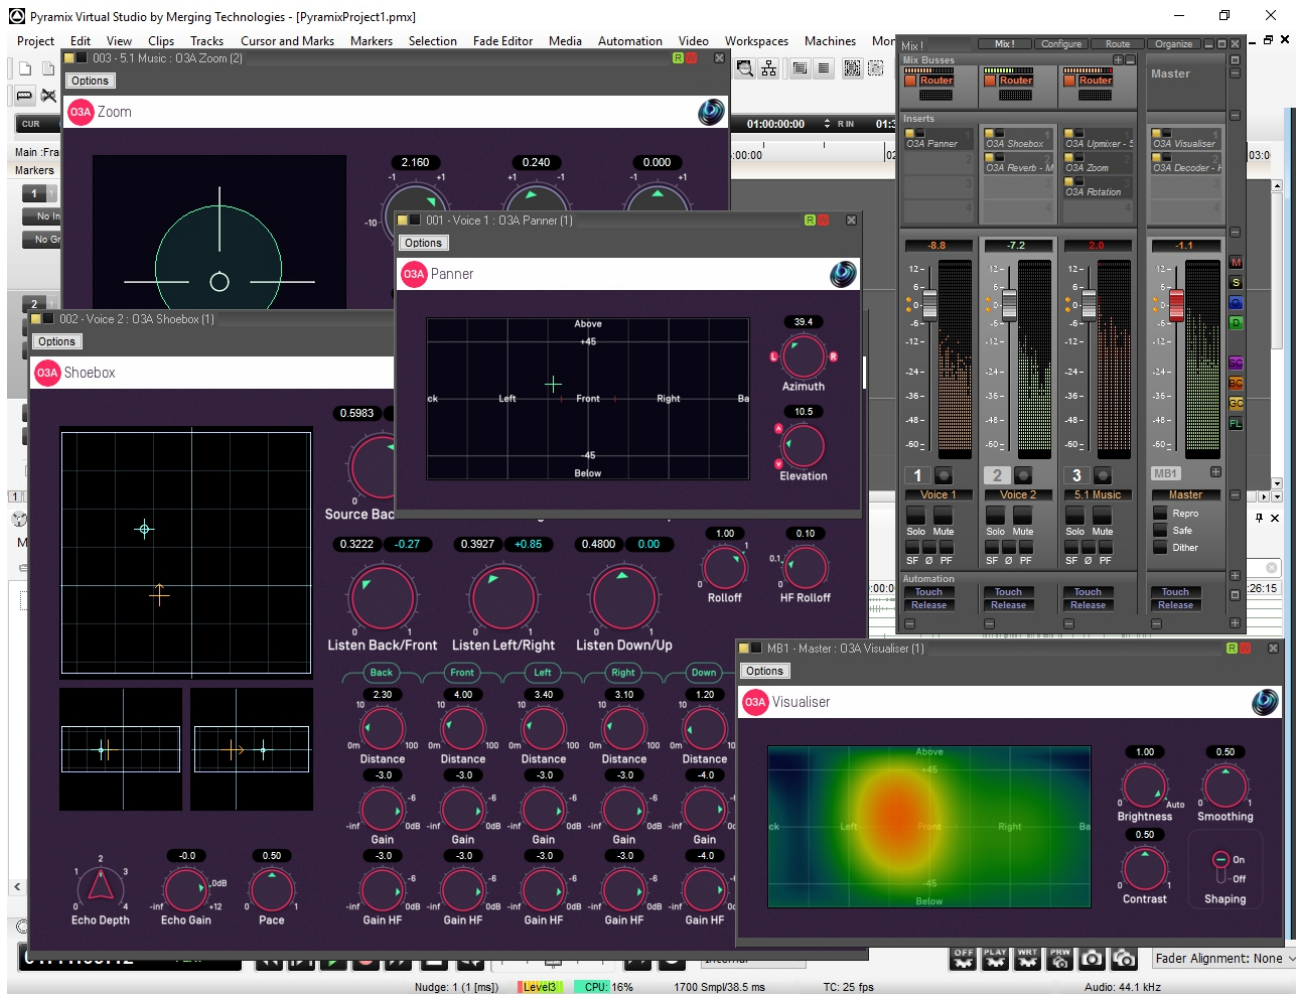


2.5.2 Mixing

Set mix busses to "Router" rather than "Pan" to ensure that the O3A plugins are determining the panning used.

In Pyramix, plugin user interface dials are set to the VST "linear" style, so drag up and down rather than around to change them.

Beyond this, plugins can be used much as you would expect!



2.6 Audio for Virtual Reality 360 Video

The **O3A streams** used by these plugins capture a spatial audio scene that can be used for Virtual Reality (VR) playback. Typically, the audio is delivered to the VR device in a raw ambisonic format and is rotated using head tracker data before live decoding. This has the effect of modifying the audio scene to take into account where the user is looking. Because raw ambisonic audio is delivered to the device, final mixes should not generally be decoded in the studio, except for monitoring.

O3A streams contain third order SN3D/ACN (AmbiX), which uses sixteen channels. This, or a subset of the channels, is directly compatible with many VR systems.

2.6.1 Rapture3D Universal

The Blue Ripple Sound Rapture3D Universal game/VR engine can use the third order SN3D audio produced by these plugins directly.

2.6.2 YouTube 360-Degree Videos

YouTube 360-degree videos can use first order SN3D. To reduce your mix to first order, just take the first four channels of your sixteen channel mix.


2.6.3 Google Jump Inspector

Google Jump Inspector can use the third order SN3D audio produced by these plugins directly.

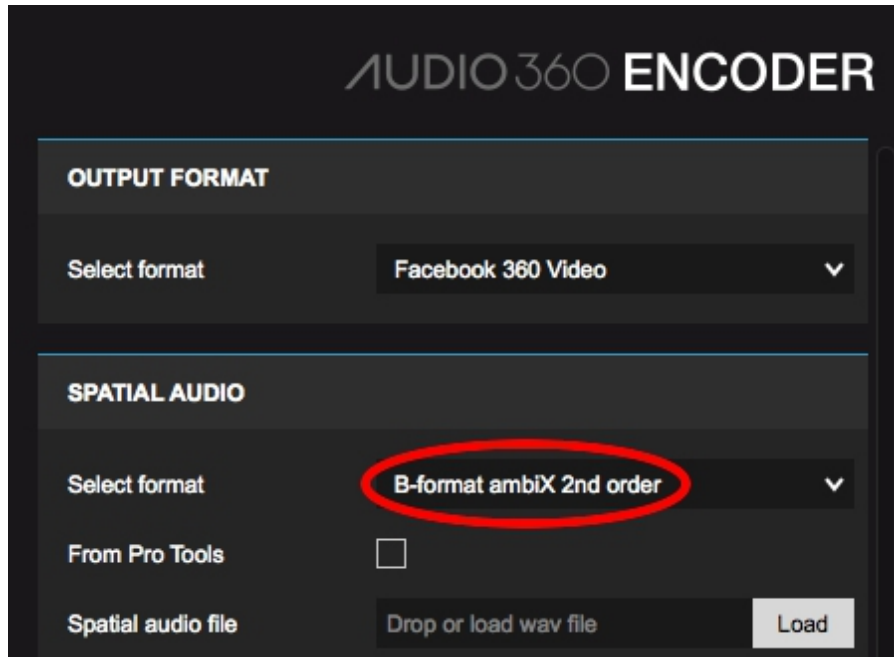
2.6.4 Facebook 360 Video

Facebook 360 videos can use second order SN3D audio.

The easiest way reduce an O3A mix to just nine channels (second order) in Reaper is often to set the render channel count to nine by *typing* the number 9 into the Channels box.



Once you have a nine channel second order file, this can be encoded using the Facebook FB360 Encoder (v1.3.0 or later). Select format "B-format ambiX 2nd order".



2.6.5 Quad Binaural

Quad binaural is an alternative system that does *not* use raw ambisonic audio. Instead, four binaural stereo images are used, corresponding to different head orientations (front, left, back and right). These are mixed together live depending on where the user is looking. This system only supports one axis of rotation, unlike ambisonics which supports all three.

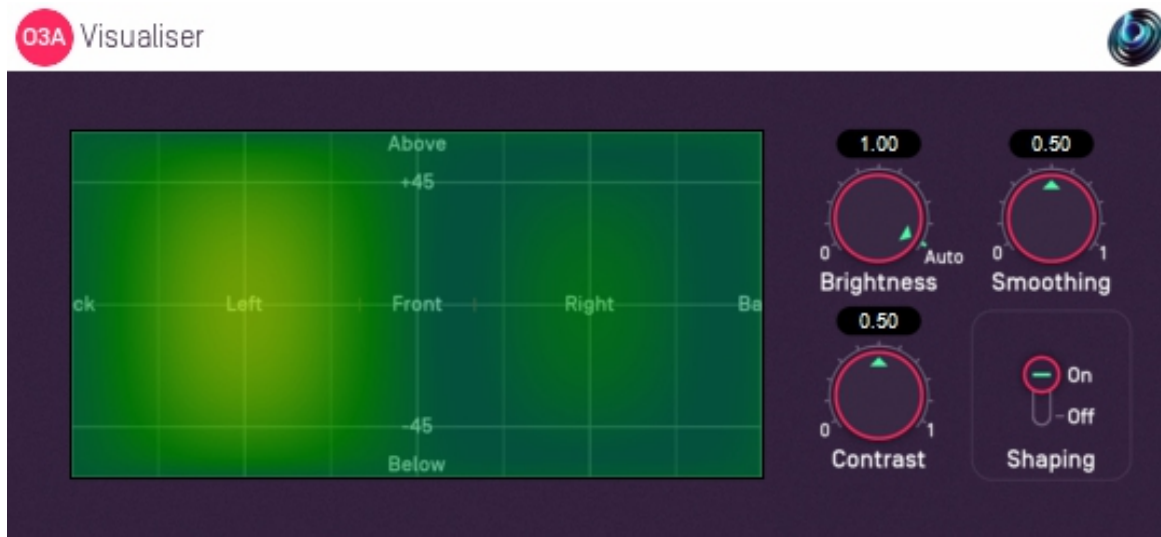
Quad binaural is generally delivered as four stereo files. To produce these, we recommend use of the the "O3A Decoder - Quad Binaural" plugin from the O3A Decoding pack. This will typically give better final results than rotating the O3A stream and using a more conventional binaural decoder to generate the Quad Binaural.

2.7 Spotting A Common Issue

In many DAWs it is easy to forget to set your track or send channel counts to sixteen at some point and end up trying to handle O3A as stereo audio. We certainly do this regularly!

This is such a common problem that it is worth taking a moment to look at what then happens, so you can spot the issue more quickly:

- Sound positioning becomes vague.
- Sound movement is only to the left and right, with nothing definite in other directions.
- The O3A Visualiser typically shows a vague left or right image, or just one colour. The left-hand case typically looks something like this:



The reason this happens is that O3A needs sixteen channels to work properly. If only the first two channels are present, these happen to be the omnidirectional and left-right ambisonic components, like in M/S stereo. These are enough to get an approximate sound level to transmit some basic left/right information, but definitely not the full detail of the soundfield!

So, keep your eyes and ears open and remember that O3A needs sixteen channels.

3 O3A Core Plugins

The O3A Core plugin library provides a number of plugin effects.

Panners:

- O3A Panner
- O3A Panner - Classic
- O3A Panner - Hemisphere
- O3A Panner - Large
- O3A Panner - Two Channel
- O3A Panner - Eight Channel

Visualisers:

- O3A Visualiser
- O3A Visualiser - Hemisphere
- O3A Colourizer
- O3A Flare

Rotation:

- O3A Look
- O3A Rotation

Decoders and converters:

- O3A Decoder - Stereo
- O3A Decoder - Mono
- O3A Decoder - 5.1 Basic
- O3A Decoder - O1A, O2A
- O3A Injector - O1A, O2A
- O3A Decoder - FuMa
- O3A Injector - FuMa

Virtual Microphones:

- O3A Beamer
- O3A Virtual Microphone

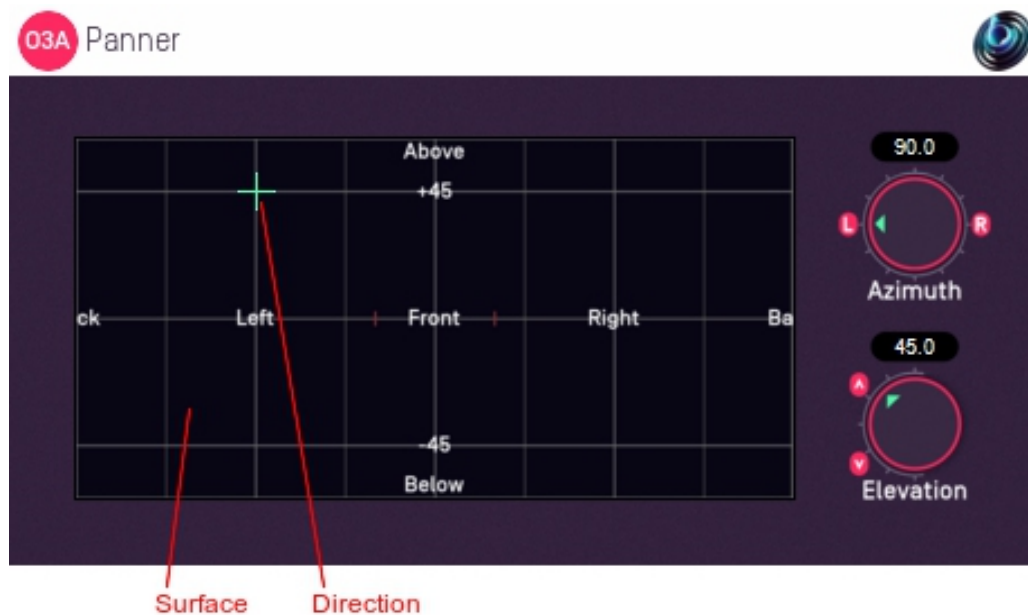
Metering:

- O3A Meter
- O3A Meter - Signal

Other:

- O3A Gain
- O3A Spatial Delay

3.1 O3A Panner



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	1	Mono
Output	16	O3A

Controls:

- [Azimuth](#)
- [Elevation](#)

3.1.1 What does this do?

This O3A panner takes a mono sound and places it in a particular direction in a 3D O3A mix.

There are various versions of this panner plugin, but this is our favourite. This version uses more screen real estate than the ["classic" version](#) but is typically easier to use, particularly if you are just starting out with this way of doing things. That said, some people prefer the ["hemisphere" version](#).

As well as the azimuth and elevation dials that the classic version also has, there is a control surface which can be used to visualise the current direction. The current direction (indicated by the crosshairs) can also be changed by clicking with the mouse. There are labels to help you find left, right, above etc. However, if this still does not

make any sense, you might want to see our page on the [O3A Visualiser](#) which lays out its viewing region in the same way.

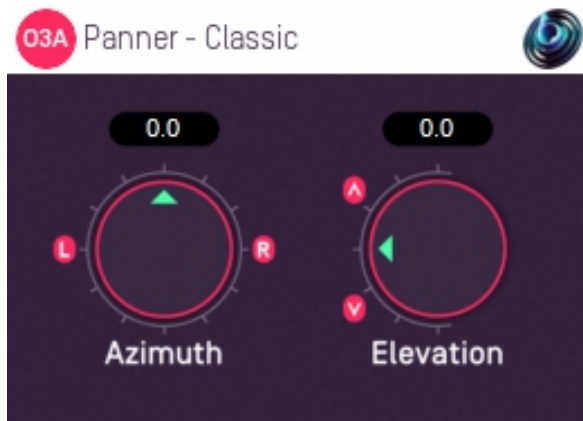
3.1.2 Controls

3.1.2.1 Azimuth and Elevation

Azimuth is the horizontal angle of the direction the source will be panned to, between -180 to +180 degrees, measured anticlockwise (left) from the front. Elevation is a vertical angle between -90 and +90 degrees, measured with positive upwards and 0 on the horizontal.

For instance, the direction for azimuth +90 and elevation +45 can be found by turning 90 degrees to the left and then looking up by 45 degrees. It may help to experiment with the control surface to get used to this way of doing things.

3.2 O3A Panner - Classic



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	1	Mono
Output	16	O3A

Controls:

- [Azimuth](#)
- [Elevation](#)

3.2.1 What does this do?

This O3A panner takes a mono sound and places it in a particular direction in a 3D O3A mix.

There are various versions of this panner plugin and this is the most "raw" version. Only the "classic" azimuth/elevation controls are provided, largely to save screen real estate. These are great if you are used to them; if you are not, you might want to try the [standard version](#).

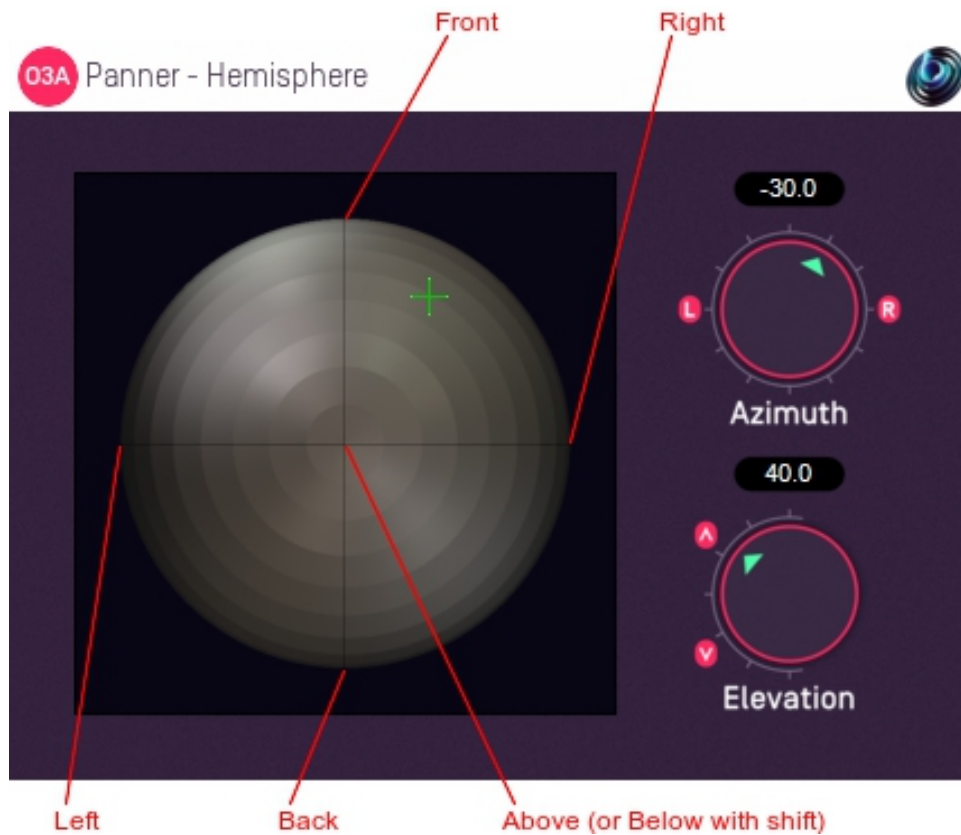
3.2.2 Controls

3.2.2.1 Azimuth and Elevation

Azimuth is the horizontal angle of the direction the source will be panned to, between -180 to +180 degrees, measured anticlockwise (left) from the front. Elevation is a vertical angle between -90 and +90 degrees, measured with positive upwards and 0 on the horizontal.

For instance, the direction for azimuth +90 and elevation +45 can be found by turning 90 degrees to the left and then looking up by 45 degrees. It may help to experiment with the [normal version](#) to get used to this way of doing things.

3.3 O3A Panner - Hemisphere



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	1	Mono
Output	16	O3A

Controls:

- Azimuth
- Elevation

3.3.1 What does this do?

This O3A panner takes a mono sound and places it in a particular direction in a 3D O3A mix.

There are various versions of this panner plugin. This version provides a control surface that maps directly onto the upper hemisphere of possible sound directions. You can click on the control surface to move the sound and holding some keys down changes how the sound is moved:

- Use the "shift" key to select the lower hemisphere (so sounds are below rather than above).
- Use the "alt" key to lock the elevation (so sounds move in a horizontal circle).

A [hemisphere visualiser](#) is available for use with this plugin. We also recommend the [standard panner](#), which has a control surface using the same rectangular layout as the standard [O3A Visualiser](#).


3.3.2 Controls

3.3.2.1 Azimuth and Elevation

Azimuth is the horizontal angle of the direction the source will be panned to, between -180 to +180 degrees, measured anticlockwise (left) from the front. Elevation is a vertical angle between -90 and +90 degrees, measured with positive upwards and 0 on the horizontal.

For instance, the direction for azimuth +90 and elevation +45 can be found by turning 90 degrees to the left and then looking up by 45 degrees. It may help to experiment with the control surface to get used to this way of doing things.

3.4 O3A Panner - Large



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	1	Mono
Output	16	O3A

Controls:

- [Azimuth](#)
- [Elevation](#)

3.4.1 What does this do?

This O3A panner takes a mono sound and places it in a particular direction in a 3D O3A mix.

This plugin is an extra large version of our [standard panner plugin](#). It can be useful on high resolution screens. It also supports background images.

3.4.2 Background Images

Press the "Load Image" button to load an image to replace the default grid background. JPG and PNG images are supported. Because of the way the plugin visualiser works, we recommend use of dark backgrounds.

The target image size is 636x316. Other sizes are crudely resized to fit.

The "Clear" button can be used to clear the image so the default grid background is shown again.

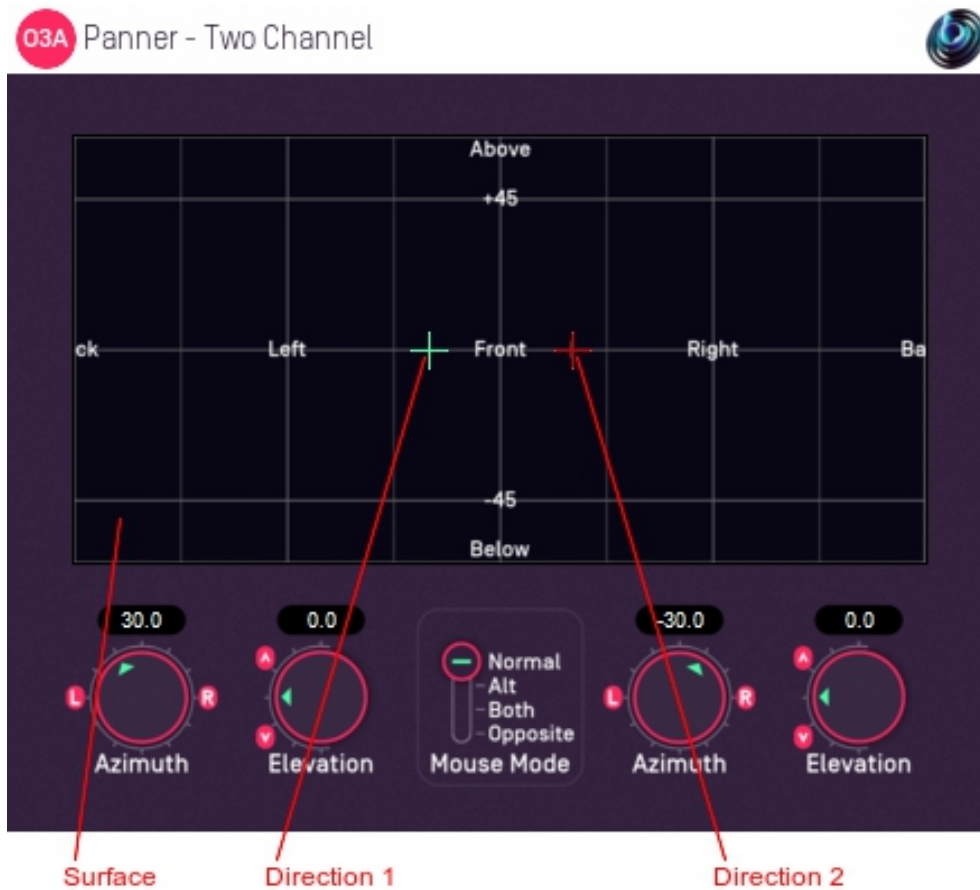
3.4.3 Controls

3.4.3.1 Azimuth and Elevation

Azimuth is the horizontal angle of the direction the source will be panned to, between -180 to +180 degrees, measured anticlockwise (left) from the front. Elevation is a vertical angle between -90 and +90 degrees, measured with positive upwards and 0 on the horizontal.

For instance, the direction for azimuth +90 and elevation +45 can be found by turning 90 degrees to the left and then looking up by 45 degrees. It may help to experiment with the control surface to get used to this way of doing things.

3.5 O3A Panner - Two Channel



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	2	Stereo
Output	16	O3A

Controls:

- Azimuth 1
- Elevation 1
- Mouse Mode
- Azimuth 2
- Elevation 2

3.5.1 What does this do?

This O3A panner takes two mono sound sources and places them in particular directions in a 3D O3A mix.

It can also be used as a simple way to upmix stereo to O3A. However, we do not particularly recommend you do this except for fairly narrow angles because a large "hole" can appear between the two sources. For a more thorough handling of material that was originally panned to stereo you wish to look at Blue Ripple Sound's "O3A Upmixer - Stereo" plugin.

There are azimuth and elevation dials for each of the two sources. Each works the same way as the controls in the [mono version](#). There is a control surface which can be used to visualise the current directions for the two sources, indicated by the green/blue and red crosshairs. These can be changed by clicking with the mouse, using the left mouse button for the first source and the right mouse button for the second. If you do not have a right mouse button, or are using a touch screen, you can hold down the 'Alt' key and press the main button instead. You can also move both sources around together using the 'Shift' key.

If pressing keys on the keyboard is inconvenient, there is a [mouse mode](#) slider on screen which can be used instead.

Mouse Mode	Key(s) Held	Function
Norm		Normal operation. Left click to move the first source, right click to move the second source.
Alt	Alt	Flips the source moved when the control surface is clicked, so a left click moves the second source rather than the first (and vice versa). Typically used to select the second source on a mouse with only one button.
Both	Shift	Changes the control surface so that movements of the mouse shift both sources.
Opp	Alt & Shift	Changes the control surface so that movements of the mouse shift both sources, but the sources move in opposite directions.

There are labels to help you find left, right, above etc. However, if this still does not make any sense, you might want to see our page on the [O3A Visualiser](#) which lays out its viewing region in the same way.

3.5.2 Controls

3.5.2.1 Azimuth 1, Elevation 1, Azimuth 2 and Elevation 2

Azimuth is the horizontal angle of the direction the source will be panned to, between -180 to +180 degrees, measured anticlockwise (left) from the front. Elevation is a vertical angle between -90 and +90 degrees, measured with positive upwards and 0 on the horizontal.

For instance, the direction for azimuth +90 and elevation +45 can be found by turning 90 degrees to the left and then looking up by 45 degrees. It may help to experiment with the control surface to get used to this way of doing things.

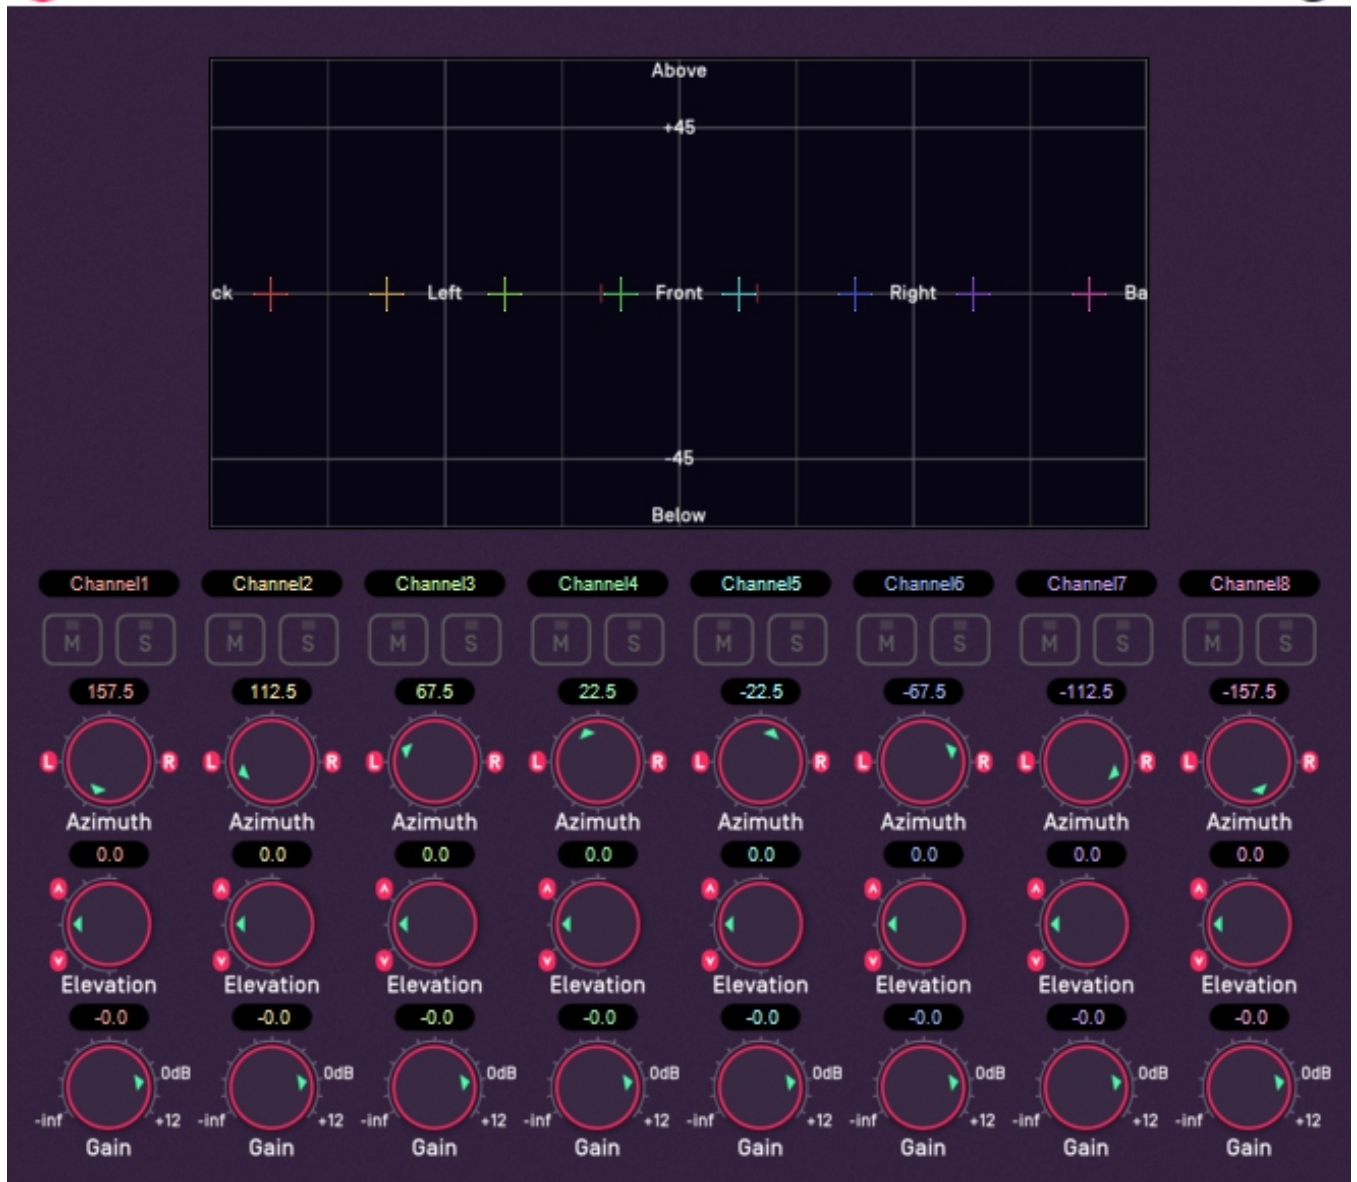
Azimuth 1 and Elevation 1 control the location used to pan the first source (which is shown in green/blue on the control surface). Azimuth 2 and Elevation 2 control the location used to pan the first source (which is shown in red).

3.5.2.2 Mouse Mode

The mouse mode slider in the middle at the bottom of the screen determine the mouse mode used when the control surface is clicked. They are an alternative to use of the Alt and Shift keys. See the [table above](#) for details on how these options affect operation.

3.6 O3A Panner - Eight Channel

O3A Panner - Eight Channel



Support:

Host Type	Support
AAX	Yes, but input is mapped to 7.1
VST2	Yes

Audio:

	Channels	Content
Input	8	Eight Channels (mapped to 7.1 for AAX)
Output	16	O3A

Controls:

- [Azimuth 1-8](#)
- [Elevation 1-8](#)
- [Gain 1-8](#)
- [Mute 1-8](#)
- [Solo 1-8](#)

3.6.1 What does this do?

This O3A panner takes eight mono sound sources and places them in particular directions in a 3D O3A mix.

There are azimuth and elevation dials for each of the sources. These work the same way as the controls in the [single channel panner](#). Each source also has mute and solo buttons and a gain control.

There is a control surface which can be used to visualise the current directions for the sources, indicated by coloured crosshairs. These can be changed by dragging with the mouse. There are labels to help you find left, right, above etc. However, if this still does not make any sense, you might want to see our page on the [O3A Visualiser](#) which lays out its viewing region in the same way.

Text can be entered into channel labels above the mute/solo buttons to help you keep track of which channel is which.

3.6.1.1 Reaper Note

If you are using this panner with Reaper, it is fairly easy to set up the individual mono source inputs as subtracks of the track on which the panner is set.

On the routing panel for each subtrack, set the "Parent Channels" to send to the relevant pair of inputs in this plugin and then pan hard left or right depending on whether this should be an odd or even input. Set this panner plugin as the first effect on the parent track and check the routing using the solo button for the source.

3.6.1.2 Pro Tools Ultimate Note

The AAX version of this plugin expects input as 7.1 audio, with the eight panning channels mapped from the 7.1 channels in the following order:

Index	Label
1	L
2	C
3	R
4	Lss
5	Rss
6	Lsr
7	Rsr
8	LFE

A simple way to use this plugin is to place it as the first plugin on a 7.1 track. Mono signals can then be sent from other tracks by addressing the individual channels of the track's 7.1 input bus. For instance, if the input bus is labelled "Panner8" then to send to the first panner channel, mono audio should be sent to "Panner8.L". To send to

the second panner channel, "Panner8.C" and so on, through to "Panner8.LFE" for the eighth channel. Routing can be checked using the plugin's solo buttons in the usual way.

3.6.2 Controls

3.6.2.1 Azimuth 1-8 and Elevation 1-8

Azimuth is the horizontal angle of the direction the source will be panned to, between -180 to +180 degrees, measured anticlockwise (left) from the front. Elevation is a vertical angle between -90 and +90 degrees, measured with positive upwards and 0 on the horizontal.

For instance, the direction for azimuth +90 and elevation +45 can be found by turning 90 degrees to the left and then looking up by 45 degrees. It may help to experiment with the control surface to get used to this way of doing things.

Azimuth and Elevation can also be changed by dragging the relevant source on the control surface.

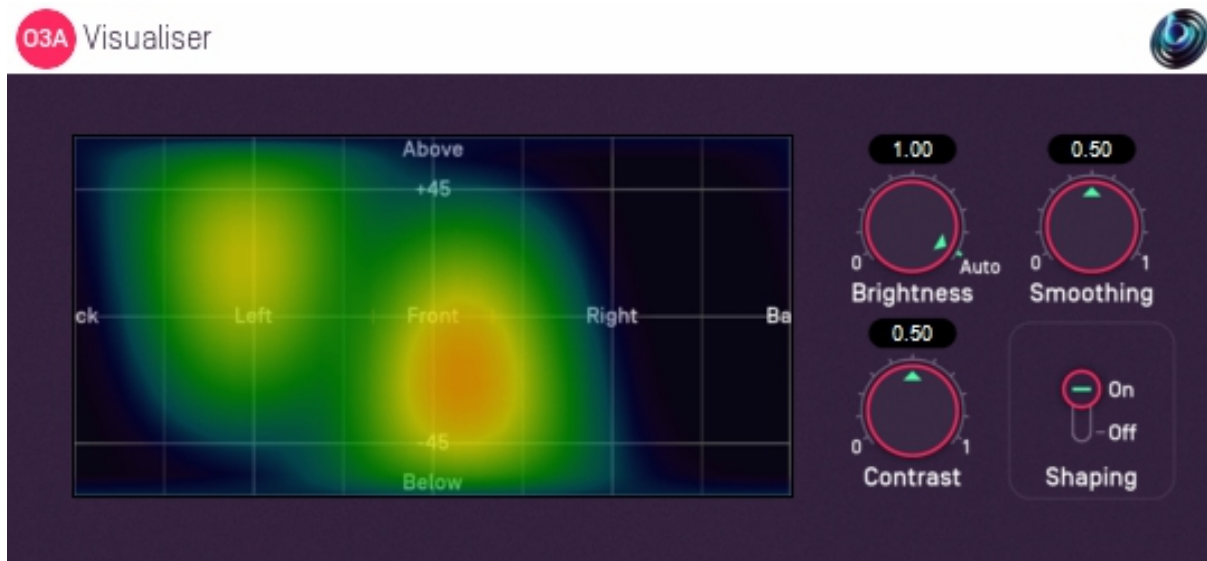
3.6.2.2 Gain 1-8

Each panned source has its own gain control, with a range from -60dB to +12dB.

3.6.2.3 Mute 1-8 and Solo 1-8

Each panned source has its own mute and solo switch. These have the usual meanings and solo overrides mute.

3.7 O3A Visualiser



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	16	O3A

Controls:

- [Brightness](#)
- [Smoothing](#)
- [Contrast](#)
- [Shaping](#)

3.7.1 What does this do?

This plugin produces a view of an O3A stream which can help understand what is happening in it. Audio is passed through unmodified, so the plugin can be inserted in the middle of an O3A effects chain.

The stream is shown using a rectangular screen region. This rectangular region shows an equal-area cylindrical projection of the directional components of the O3A soundfield interpreted over a sphere.

Or, to put it another way, it lights up in different places depending on where a sound comes from. Some of those places are labelled in the image. The view is like a map of the world, where sounds above are at the "north pole" and are shown across the top of the view and sounds below are at the "south pole" at the bottom. Sounds rotating

around at head height move around the equator, so they move along the view halfway down, wrapping around between the left and right edges as they pass behind.

A good way to try this out is to chain the [O3A Panner](#) immediately before this plugin as this uses the same directional approach. Then set up a decoder after this suitable for use with your speakers or headphones. Then you can play around with different sound directions and hear *and see* the results.

To reduce CPU load, this visualiser stops most processing when its user interface is not on screen.

3.7.2 Controls

3.7.2.1 Brightness

This controls the brightness of the image. If the screen is black when sounds are playing, you may wish to turn this up. If it is all red, turn it down.

Alternatively, you can set the dial to "Auto", in which case the plugin will try to ride the brightness control to keep the visualisation useful.

3.7.2.2 Smoothing

The image generated is smoothed over time. Use more smoothing for a more stable, slower moving image. Use less for a more responsive one.

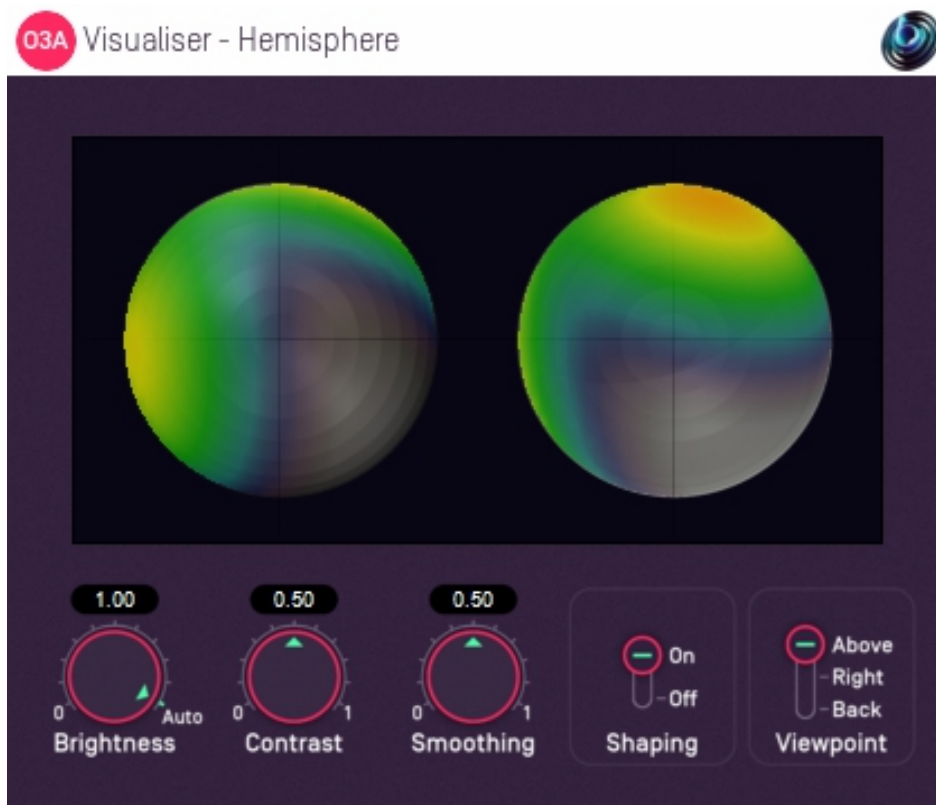
3.7.2.3 Contrast

This controls the contrast of the image. Higher values produce a sharper image but map a smaller dynamic range onto the screen.

3.7.2.4 Shaping

Shaping should normally be left on as it produces a more straightforward view. If this is turned off, the image will typically be sharper, but ripples will be visible in directions away from original sound directions.

3.8 O3A Visualiser - Hemisphere



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	16	O3A

Controls:

- Brightness
- Smoothing
- Contrast
- Shaping
- Viewpoint

3.8.1 What does this do?

This plugin produces a view of an O3A stream which can help understand what is happening in it. Audio is passed through unmodified, so the plugin can be inserted in the middle of an O3A effects chain.

The stream is shown using a two circular regions each showing a projection of a half of the directional components of the O3A soundfield interpreted over a sphere.

Or, to put it another way, it lights up in different places depending on where a sound comes from.

By default, the left sphere shows the upper hemisphere of the soundfield and the right sphere shows the lower hemisphere, as if both were viewed from above. The image is orientated so that sounds at the front appear at the top of both images. However, the [viewpoint](#) can be changed.

A good way to try this out is to chain the [O3A Panner - Hemisphere](#) plugin immediately before this one as this uses the same directional approach (as long as the [viewpoint](#) is set to "Above"). Then set up a decoder after this suitable for use with your speakers or headphones. Then you can play around with different sound directions and hear and see the results.

We do not particularly recommend this visualisation. Generally, we suggest you use the standard [O3A Visualiser](#) instead.

To reduce CPU load, this visualiser stops most processing when its user interface is not on screen.

3.8.2 Controls

3.8.2.1 Brightness

This controls the brightness of the image. If the visualisation shows nothing when sounds are playing, you may wish to turn this up. If it is all red, turn it down.

Alternatively, you can set the dial to "Auto", in which case the plugin will try to ride the brightness control to keep the visualisation useful.

3.8.2.2 Smoothing

The image generated is smoothed over time. Use more smoothing for a more stable, slower moving image. Use less for a more responsive one.

3.8.2.3 Contrast

This controls the contrast of the image. Higher values produce a sharper image but map a smaller dynamic range onto the screen.

3.8.2.4 Shaping


Shaping should normally be left on as it produces a more straightforward view. If this is turned off, the image will typically be sharper, but ripples will be visible in directions away from original sound directions.

3.8.2.5 Viewpoint

By default, the visualisation shows the upper hemisphere on the left and the lower hemisphere on the right, both viewed as the viewpoint was from **above**.

Alternatively, the viewpoint can be changed to view from the **right** or from **behind**. In both cases sounds above appear at the top of the image.

3.9 O3A Colourizer



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	16	O3A

Controls:

- Hue Cutoff
- Smooth Cutoff
- Rider Cutoff
- Brightness

3.9.1 What does this do?

This plugin is a variant of the [O3A Visualiser](#) plugin. Like that plugin, it produces a view of an O3A stream which can help understand what is happening in it. Audio is passed through unmodified, so the plugin can be inserted in the middle of an O3A effects chain.

The stream is shown using a rectangular screen region. This rectangular region shows an equal-area cylindrical projection of the directional components of the O3A soundfield interpreted over a sphere.

Or, to put it another way, it lights up in different places depending on where a sound comes from. Some of those places are labelled in the image. The view is like a map of the world, where sounds above are at the "north pole" and are shown across the top of the view and sounds below are at the "south pole" at the bottom. Sounds rotating around at head height move around the equator, so they move along the view halfway down, wrapping around between the left and right edges as they pass behind.

A good way to try this out is to chain the [O3A Panner](#) immediately before this plugin as this uses the same directional approach. Then set up a decoder after this suitable for use with your speakers or headphones. Then you can play around with different sound directions and hear *and see* the results.

The main differences between this and the [O3A Visualiser](#) are the size of the plugin and how colours are mapped. With this plugin, colours are chosen using a hue that depends on the frequency content of the sound. Sounds that are predominantly low-frequency are coloured red and sounds that are predominantly high-frequency are coloured magenta. Sounds between are coloured with other hues accordingly.

To reduce CPU load, this visualiser stops most processing when its user interface is not on screen.

3.9.2 Background Images

Press the "Load Image" button to load an image to replace the default grid background. JPG and PNG images are supported. Because of the way the plugin visualiser works, we recommend use of dark backgrounds.

The target image size is 636x316. Other sizes are crudely resized to fit.

The "Clear" button can be used to clear the image so the default grid background is shown again.

3.9.3 Controls

3.9.3.1 Hue Cutoff

The plugin selects hues by splitting the sound into low and high frequencies. To do this, it uses a low-pass filter with a cutoff frequency given by this control. Generally, it is best to set this where you want the greatest sensitivity to frequency content.

3.9.3.2 Smooth Cutoff

The local image level is smoothed over time using a low-pass filter with a cutoff frequency given by this control. Lower values will give a more stable, slower moving image. Higher values will give a more responsive one.

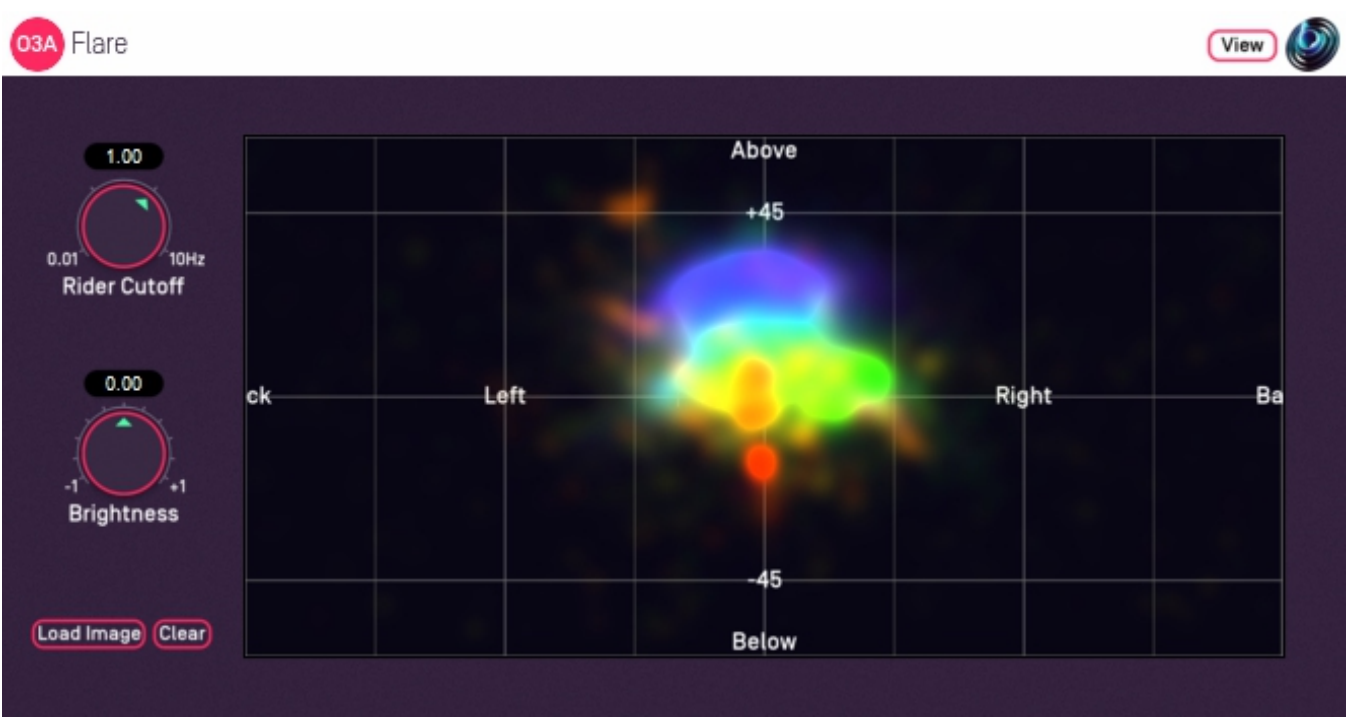
3.9.3.3 Rider Cutoff

A "gain rider" tracks the overall signal level using another low-pass filter. Lower values respond more slowly to changes in the overall level of the signal. Higher values respond more quickly.

3.9.3.4 Brightness

This modifies the overall brightness of the image.

3.10 O3A Flare



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	16	O3A

Controls:

- [Rider Cutoff](#)
- [Brightness](#)

3.10.1 What does this do?

This plugin is a variant of the [O3A Visualiser](#) plugin. Like that plugin, it produces a view of an O3A stream which can help understand what is happening in it. Audio is passed through unmodified, so the plugin can be inserted in the middle of an O3A effects chain.

The stream is shown using a rectangular screen region. This rectangular region shows an equal-area cylindrical projection of the directional components of the O3A soundfield interpreted over a sphere.

Or, to put it another way, it lights up in different places depending on where a sound comes from. Some of those places are labelled in the image. The view is like a map of the world, where sounds above are at the "north pole" and are shown across the top of the view and sounds below are at the "south pole" at the bottom. Sounds rotating around at head height move around the equator, so they move along the view halfway down, wrapping around between the left and right edges as they pass behind.

A good way to try this out is to chain the [O3A Panner](#) immediately before this plugin as this uses the same directional approach. Then set up a decoder after this suitable for use with your speakers or headphones. Then you can play around with different sound directions and hear *and see* the results.

This plugin uses substantially different processing to the [O3A Visualiser](#) and [O3A Colourizer](#). This plugin uses simple first order estimates of the predominant sound direction in a range of frequency bands. These directions are painted using colour coding, ranging from red for low frequencies up to magenta for high frequencies (up to about 20kHz). Plot accuracy is better for simpler scenes; multiple sources with similar frequency content will often result in smudged or misaligned images.

To reduce CPU load, this visualiser stops most processing when its user interface is not on screen.

3.10.2 Background Images

Press the "Load Image" button to load an image to replace the default grid background. JPG and PNG images are supported. Because of the way the plugin visualiser works, we recommend use of dark backgrounds.

The target image size is 512x256. Other sizes are crudely resized to fit.

The "Clear" button can be used to clear the image so the default grid background is shown again.

3.10.3 Controls

3.10.3.1 Rider Cutoff

A "gain rider" tracks the overall signal level using another low-pass filter. Lower values respond more slowly to changes in the overall level of the signal. Higher values respond more quickly.

3.10.3.2 Brightness

This modifies the overall brightness of the image.

3.11 O3A Look



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	16	O3A

Controls:

- [Azimuth](#)
- [Elevation](#)

3.11.1 What does this do?

This plugin changes where "front" is in the 3D audio scene, as if the listener had turned their head to look (and listen) in another direction.

The "look" direction is chosen using the dials or control surface, in the normal way (for instance, see the [O3A Panner](#)). Sounds that were in this direction will be to the front after processing.

This is essentially an simplified (and inverted) version of the [O3A Rotation](#) plugin. The "roll" and "ordering" controls are not provided. Instead, we just assume that the listener turns their head and pitches it backwards or forwards. Rolling to the sides is not supported (use the rotation plugin if you need this).

3.11.2 Controls

3.11.2.1 Azimuth and Elevation

Azimuth is the horizontal angle of the "look" direction, between -180 to +180 degrees, measured anticlockwise (left) from the front. Elevation is a vertical angle between -90 and +90 degrees, measured with positive upwards and 0 on the horizontal.

3.12 O3A Rotation



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	16	O3A

Controls:

- Yaw
- Pitch
- Roll
- Ordering

3.12.1 What does this do?

This plugin rotates the entire O3A soundfield around the centre of the room.

Any rotation around the centre is possible and is built up using yaw, pitch and roll rotations. You can think of the rotations as being applied in order.

We *strongly recommend* that you use the [O3A Visualiser](#) when working with this plugin, particularly if you are going to use more than one of the dials at once. Otherwise, this plugin can be very confusing!

If you are using this plugin to simulate changes to the listener's orientation, you might also want to consider the [O3A Look](#) plugin, which is simpler to use but does not include a "roll" control.

3.12.2 Controls

3.12.2.1 Yaw, Pitch and Roll

These controls use angles between -180 and +180 degrees which can each be thought of as controlling individual rotations around three fixed axes. The overall rotation is a combination of the three.

Rotation	Axis	Description of Positive Movement
Yaw	Z	The soundfield is "rotated" left around the middle of the room. For instance, a sound that starts at the front moves left, a sound that starts at the left moves backwards and a sound directly above does not move.
Pitch	Y	The soundfield is "tumbled" backwards. For instance, a sound that starts at the front moves upwards, a sound that starts at the left does not move and a sound directly above moves backwards.
Roll	X	The soundfield is "tilted" clockwise. For instance, a sound that starts at the front does not move, a sound that starts at the left moves upwards and a sound directly above moves right.

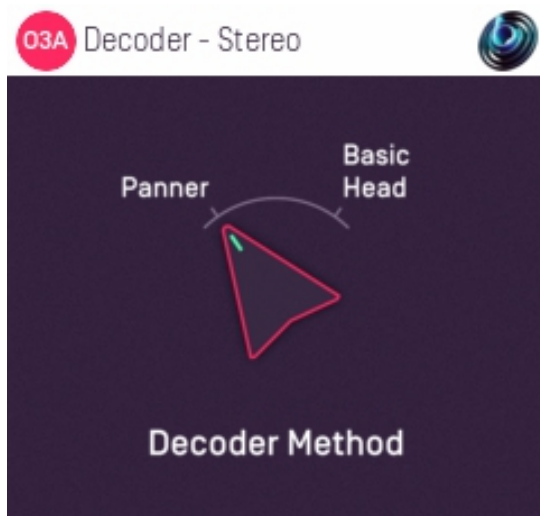
This plugin can be used to simulate changes in the listener's orientation. If you are doing this, you will probably want to change the sign of the angles involved. For instance, turning the listener's head to the right means that the listener hears sound move to the left.

3.12.2.2 Ordering

This control determines the order in which the yaw, pitch and roll rotations are applied to make up a single rotation. The default is Y-P-R (i.e. yaw, pitch and then roll) but this can be switched to R-P-Y (i.e. roll, pitch and then yaw).

When simulating changes in the listener's orientation, we generally recommend the Y-P-R ordering. Then, the three dials can be thought of as turning the listener's head horizontally, then pitching it forwards or backwards and finally rolling it from side to side.

3.13 O3A Decoder - Stereo



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	2	Stereo

Controls:

- [Decoder Method](#)

3.13.1 What does this do?

This plugin takes a third order ambisonics (O3A) mix and decodes it as a simple, robust stereo mix.

3.13.2 Controls

3.13.2.1 Decoder Method

Two decoder methods are supported, "Panner" and "Basic Head".

Method	Description
Panner	<p>This method produces a simple stereo mix, largely as if the sounds within the mix had been panned conventionally. This can work well on speakers or headphones.</p> <p>Sounds from the front and back are folded down on the stereo image using different angle conventions. Specifically, sounds panned from -30 to +30 degrees at the front will pass across the stereo image, as will sounds panned from -110 to +110 degrees at the back. Sounds above or below are fed to both speakers.</p>
Basic Head	<p>This method produces a stereo mix using a very crude model of a human head. When listened to on headphones, this imitates typical left and right ear levels that would occur naturally, albeit not in a frequency-dependent way. This can work well on speakers or headphones.</p> <p>This is not a binaural decoder. The data for this method is based on the Amber HRTF headphone decoder available in the Blue Ripple Sound O3A Decoding library and Rapture3D, but massively simplified so that the output is not frequency-dependent. This reduces the strength of 3D cues on headphones, but means that the frequency content is changed less when played on stereo speakers. For full binaural decoding, we recommend the O3A Decoding or O3A View plugin libraries, or Rapture3D Advanced.</p> <p>Unlike most other O3A decoders, this decoder presents sounds from different directions at different levels. For instance, sounds to the front are louder than those behind.</p>

As an alternative to this plugin, you can use the "[O3A Virtual Microphone](#)" plugin to produce a stereo mix using a virtual stereo microphone. Additionally, the "O3A Decoder - Headphones" plugin from the O3A Decoding plugin library can be used to synthesize binaural stereo.

3.13.3 Technical Notes

The Blue Ripple Sound Amber HRTF uses data from the IRCAM LISTEN HRTF data set, available at <http://recherche.ircam.fr/equipes/salles/listen/index.html>. It has been processed and simplified heavily.

However, please again note that **this is not a binaural decoder**.

3.14 O3A Decoder - Mono



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	1	Mono

3.14.1 What does this do?

This plugin takes a third order ambisonics (O3A) mix and reduces it to mono by discarding everything but its omnidirectional component.

This type of mix is particularly suitable for level analysis, including loudness monitoring.

If you wish to extract a mono response from the soundfield in a particular direction, you might also want to try the [O3A Virtual Microphone](#) or [O3A Beamer](#) plugins.

3.15 O3A Decoder - 5.1 Basic



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	6	5.1 Surround

3.15.1 What does this do?

This plugin takes a third order ambisonics (O3A) mix and decodes it as a 5.1 surround mix.

A resilient decoding is used, suitable for small or large venues. ITU speaker angles are assumed.


3.15.2 Channels

The channel ordering for the 5.1 surround mix is:

Channel	AAX	VST
1	Front Left	Front Left
2	Front Centre	Front Right
3	Front Right	Front Centre
4	Left Surround	LFE
5	Right Surround	Side Left
6	LFE	Side Right

The low frequency effect channel (the ".1" of the 5.1 format) is *not* used. Sounds above or below are fed to all speakers (except the LFE).

3.16 O3A Decoder - O1A, O2A



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

O3A Decoder - O1A	Channels	Content
Input	16	O3A
Output	4	O1A

O3A Decoder - O2A	Channels	Content
Input	16	O3A
Output	9	O2A

3.16.1 What does this do?


These simple plugins take a third order (O3A) ambisonic mix and convert it to a lower resolution first order (O1A) or second order (O2A) mix. In all cases, the SN3D convention is used.

These plugins are extremely simple and just reduce the channel count in use by discarding the later channels, which has the effect of reducing the spatial detail of the mix.

Note that some DAWs (such as Reaper) pass through channels that are not written to. These plugins will appear to have no effect in these DAWs!

These plugins can also be "misused" to reduce the order of N3D/ACN or FuMa material. This is a rare case where the same plugin will work correctly for all three ambisonic formats.

3.17 O3A Injector - O1A, O2A



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

O3A Injector - O1A	Channels	Content
Input	4	O1A
Output	16	O3A

O3A Injector - O2A	Channels	Content
Input	9	O2A
Output	16	O3A

3.17.1 What does this do?


These simple plugins take a lower resolution first order (O1A) or second order (O2A) mix and convert it to a third order (O3A) ambisonic mix. In all cases, the SN3D convention is used.

These plugins are extremely simple and just increase the channel count in use, filling new channels with silence. Although this increases the order of the mix, it does not increase its spatial detail. Further, the resulting mix is likely to have somewhat unnatural features acoustically.

More sophisticated means to convert first order (O1A) to third order include the "O3A Injector - First Order" plugin from the O3A Upmixers pack and the "O3A Harpex Upsampler" plugin from the O3A Harpex Upsampler pack.

These plugins can also be "misused" to increase the order of N3D/ACN or FuMa material. This is a rare case where the same plugin will work correctly for all three ambisonic formats.

3.18 O3A Decoder - FuMa



Support:

Host Type	Support
AAX	Yes, but output is mapped to a standard ambisonic (SN3D) stem
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	16	FuMa Ambisonic Audio (mapped to standard ambisonic stems for AAX)

3.18.1 What does this do?

This plugin takes a third order ambisonic mix prepared with these plugins (O3A) and converts it to a third order mix using the FuMa ambisonic convention. The FuMa convention was used by versions of these Blue Ripple Sound plugins prior to version 2.0 and is an extension of classic First Order Ambisonics (using "WXYZ").

We have labelled this plugin as a decoder hopefully to make it clear that its output is an external format and no longer compatible with the [convention that the rest of these plugins use](#).

You can convert back using the [O3A Injector - FuMa](#) plugin.


In Pro Tools Ultimate, as the FuMa format is not supported directly as a stem format, this plugin outputs audio using the ordinary ambisonic stem format. Normally this should be encoded as O3A (i.e. SN3D/ACN), so we are "misusing" the format here.

3.18.2 Orders

This plugin produces third order FuMa output. If you only want first or second order output, just take the first four or nine output channels and ignore the rest.

Alternatively, the [O3A Decoder - O1A or O2A plugins](#) can be "misused" to reduce the order of the FuMa material as this is a rare case where a plugin will work correctly for both SN3D and FuMa.

3.19 O3A Injector - FuMa



Support:

Host Type	Support
AAX	Yes, but input is mapped to a standard ambisonic (SN3D) stem
VST2	Yes

Audio:

	Channels	Content
Input	16	FuMa Ambisonic Audio (mapped to standard ambisonic stems for AAX)
Output	16	O3A

3.19.1 What does this do?

This plugin takes a mix in the FuMa ambisonic convention and converts it to [third order ambisonics as used by the current version of these plugins \(O3A\)](#). The FuMa convention was used by versions of these Blue Ripple Sound plugins prior to version 2.0.

FuMa is an extension of classic First Order Ambisonics (using just the WXYZ channels) and material in this format can also be converted using this plugin.

You can convert back with the [O3A Decoder - FuMa](#) plugin.

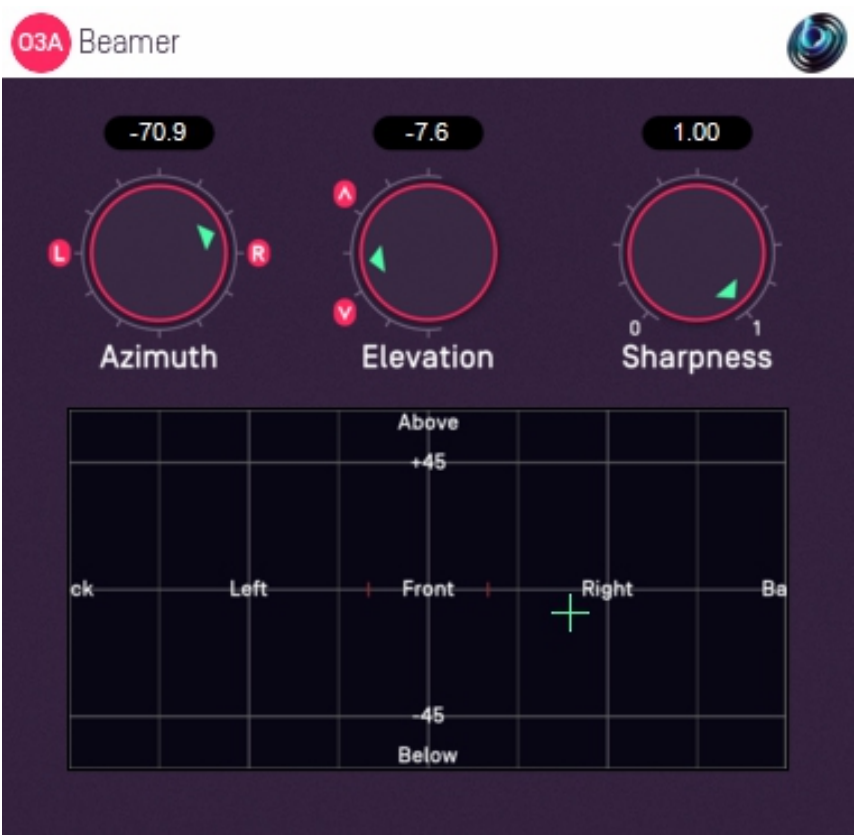
In Pro Tools Ultimate, as the FuMa format is not supported directly as a stem format, this plugin expects audio using the ordinary ambisonic stem format. Normally this should be encoded as O3A (i.e. SN3D/ACN), so we are "misusing" the format here.

3.19.2 Orders

This plugin expects third order FuMa input. If you only have first or second order, feed this to the first four or nine input channels and feed silence to the rest.

Alternatively, the [O3A Injector - O1A or O2A plugins](#) can be "misused" to increase the order of the FuMa material as this is a rare case where a plugin will work correctly for both SN3D and FuMa.

3.20 O3A Beamer



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	1	Mono

Controls:

- Azimuth
- Elevation
- Sharpness

3.20.1 What does this do?

The O3A Beamer extracts sound in a particular "beam" direction from a 3D O3A mix. You can think of this as a narrow directional microphone response, recording in a particular direction in the soundfield.

The beam direction can be set using the azimuth and elevation controls, or the control surface, where the direction is indicated by the crosshairs. There are labels to help you find left, right, above etc. However, if this still does not make any sense, you might want to see our page on the [O3A Visualiser](#) which lays out its viewing region in the same way.

The algorithm used is a "passive" one and should not colour the frequency content of the sound significantly.

If you are interested in producing virtual microphone responses from an O3A mix, you may also wish to read about the [O3A Virtual Microphone](#).

3.20.2 Controls

3.20.2.1 Azimuth and Elevation

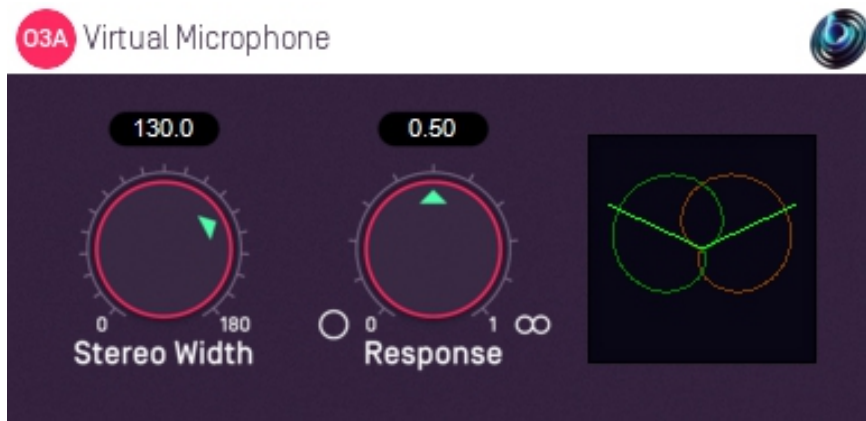
Azimuth is the horizontal angle of the direction to be extracted, between -180 to +180 degrees, measured anticlockwise (left) from the front. Elevation is a vertical angle between -90 and +90 degrees, measured with positive upwards and 0 on the horizontal.

For instance, the direction for azimuth +90 and elevation +45 can be found by turning 90 degrees to the left and then looking up by 45 degrees. It may help to experiment with the [O3A Panner](#) to get used to this way of doing things.

3.20.2.2 Sharpness

The sharpness control determines how directional the beam response is. A value of zero reduces the response to an omnidirectional one like that of the [mono decoder](#). A value of one is the usual case and gives the sharpest response.

3.21 O3A Virtual Microphone



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	2	Stereo

Controls:

- Stereo Width
- Response

3.21.1 What does this do?

The O3A Virtual Microphone extracts a stereo image from a 3D O3A mix using a front-facing "virtual" stereo microphone placed at the centre of the soundfield. The angle between the two capsules of this virtual microphone can be varied, along with their directional response. Available directional responses include omni, cardioid and figure-of-eight. This makes it possible to synthesise the results of various conventional coincident stereo recording techniques such as Blumlein X-Y and M/S.

This plugin does not allow you to point the front of the virtual microphone in other directions. If the front is not where you want, remember that you can rotate the soundfield using the [O3A Rotation](#) plugin before using this one.

This plugin actually only needs first order ambisonic information (and even then, the height channel is not used).

If you are interested in extracting sounds in particular directions from an O3A mix, you may also be interested in the sharper response of the [O3A Beamer](#) plugin.

3.21.2 Controls

3.21.2.1 Stereo Width

This is the angle between the two virtual microphones, between 0 and 180 degrees. For instance, 90 degrees should be used for Blumlein X-Y.

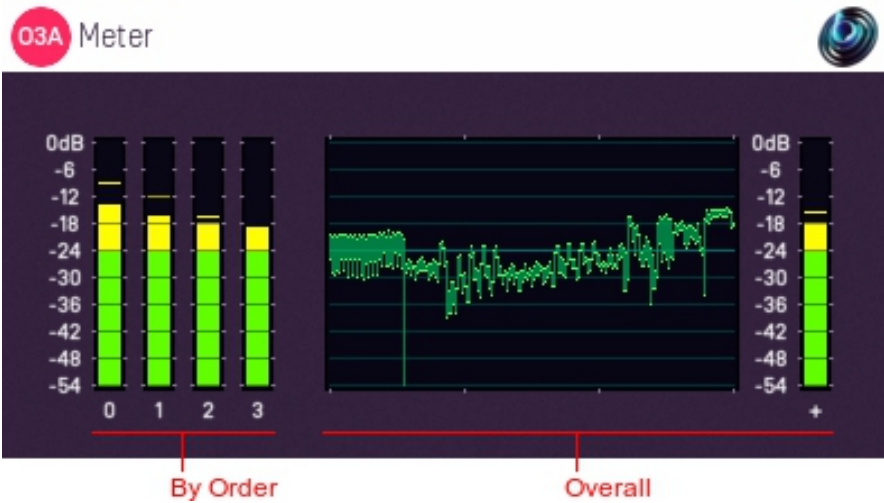
3.21.2.2 Response

The response control determines the directional response of the two microphone capsules. Some useful values are:

Response Value	Name
0.00	Omnidirectional
0.25	Subcardioid
0.50	Cardioid
0.75	Hypercardioid
1.00	Figure-of-Eight

For instance, 1.00 (figure-of-eight) should be used for Blumlein X-Y.

3.22 O3A Meter



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	16	O3A

3.22.1 What does this do?

This plugin measures the sound level in an O3A stream. A breakdown by ambisonic order is provided, along with a history graph and the current overall level. Audio is passed through unchanged.

3.22.2 Meters

From left to right, the following displays are shown:


1. Bar meter showing the current and recent peak level in order 0 (ACN channel 0 only).
2. Bar meter showing the current and recent peak level in order 1 (ACN channels 1 to 3).
3. Bar meter showing the current and recent peak level in order 2 (ACN channels 4 to 8).
4. Bar meter showing the current and recent peak level in order 3 (ACN channels 9 to 15).
5. History min/max graph of overall level (all 16 channels).
6. Bar meter showing the overall current and peak level (all 16 channels).

Each bar meter shows the current RMS level in green, averaged using a 400ms rectangular window and translated to a decibel scale, along with a recent peak calculation using a 2s memory. The history graph shows the last three minutes of the overall level, displaying the range from minimum to maximum for each second.

The meters are all normalised so that a 0dBFS peak sine wave panned in any direction will produce a -3.01dB output. There is no sensitivity to frequency or direction. If you are interested in measuring loudness, you may wish to consider the "O3A Meter - Karma" plugin from the [Blue Ripple Sound](#) O3A Decoding library, which will produce comparable numbers when set to use LUFS units, subject to directional weighting.

For material that has been assembled by simple panning, the five bar meters should be roughly in line with each other.

3.23 O3A Meter - Signal



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	16	O3A

3.23.1 What does this do?


This extremely simple plugin measures the signal peak in each of the 16 channels of an [O3A stream](#) to indicate if a signal is present. It is intended for basic diagnostics (e.g. "are all my channels getting through?") and not for serious signal analysis. Audio is passed through unchanged.

3.23.2 Meters

Sixteen lights are shown, for each of the sixteen channels in O3A. They are shown in rows corresponding to the ambisonic orders. Lights switch on with a yellow colour in the presence of very low signal levels (-100dB and above) and switch to red at high levels (0dB and above). These signal levels are simple peak levels, *not* "True Peak".

Note that first order ambisonic B-Format only uses the first four channels of O3A, in which case only the first four lights (top two rows) will be on. Sounds exactly on the horizontal plane switch on alternating lights on each row and a few special directions do other strange things. However, normal 3D streams should switch on all sixteen lights; if they do not there is probably something wrong! A common error is to forget to set a track's channel count or send to sixteen or more and this plugin will usually make that obvious.

3.24 O3A Gain



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	16	O3A

Controls:

- [Gain](#)

3.24.1 What does this do?

This trivial plugin applies a gain to the O3A stream. All channels of the stream are affected equally.

3.24.2 Controls

3.24.2.1 Gain

The gain to apply, in decibels, between -30dB and +30dB.

3.25 O3A Spatial Delay



Support:

Host Type	Support
AAX	Yes
VST2	Yes

Audio:

	Channels	Content
Input	16	O3A
Output	16	O3A

Controls:

- Delay Type
- Delay Time
- Dry Gain
- Wet Gain
- Feedback Gain
- LPF Cutoff
- Reflection Type
- Yaw
- Pitch
- Roll
- Ordering
- Clear Delay

3.25.1 What does this do?

This plugin delays a third order ambisonic (O3A) stream, applying a scene reflection, rotation and low pass filter to the delayed sound. It also supports feedback, so the output of the plugin can be fed back to its input to create repeating sound effects.

The rotation module of this plugin operates in the same way as the [O3A Rotation](#) plugin and can be controlled by an external View application. The reflection, rotation and low pass filter appear *within* the feedback loop. The wet gain does not, applying only to the final output.

The "Clear Delay" button on the bottom right can be used to clear the audio in the delay line.

3.25.2 Controls

3.25.2.1 Delay Type and Delay Time

The delay type determines whether the delay time is specified in beats or in seconds. When beats are used, the delay length is a multiple of the beat length in your DAW's current tempo, so if your tempo is in crotchets, then setting the delay to 2.0 results in a delay time equivalent to a minim.

Delay times are rounded to the nearest sample and the minimum delay is one sample. The maximum setting is 20 seconds or 20 beats.

Changes to the delay time are implemented using fades between different delay settings. This means that the delay cannot be used to simulate Doppler shift, but that changes are less likely to be disruptive to musical material. Continuous changes to the delay time may result in artefacts.

Please note that high delay time settings use a significant amount of memory (RAM). This memory is not released until the plugin instance is unloaded.

3.25.2.2 Dry and Wet Gain

These control how much unmodified "dry" input and delayed "wet" material is included in the output.

3.25.2.3 Feedback Gain

The feedback gain determines how much of the delayed output is sent back into the delay line. The wet gain is not applied here, but rotation and the low pass filter are.

When the feedback gain is set at or close to 0dB, the delay output level is liable to increase over time. Please be careful not to damage anything because of this. The delay line can be cleared using the Clear button, but being conservative with the feedback gain setting is normally recommended.

3.25.2.4 Reflection Type

This can be used to apply a fixed spatial reflection to the O3A material before a spatial rotation is applied. The spatial reflections available are:

- None
- Left/Right
- Front/Back
- Up/Down

Note that, by combination with the rotation, it is possible to construct any other angle of reflection.

3.25.2.5 Yaw, Pitch and Roll

A scene rotation is applied to the delayed O3A scene, in a similar way to the rotation used by the [O3A Rotation](#) plugin. Feedback can result in the rotation being applied more than once to the same material.

These controls use angles between -180 and +180 degrees which can each be thought of as controlling individual rotations around three fixed axes. The overall rotation is a combination of the three.

Rotation	Axis	Description of Positive Movement
Yaw	Z	The soundfield is "rotated" left around the middle of the room. For instance, a sound that starts at the front moves left, a sound that starts at the left moves backwards and a sound directly above does not move.
Pitch	Y	The soundfield is "tumbled" backwards. For instance, a sound that starts at the front moves upwards, a sound that starts at the left does not move and a sound directly above moves backwards.
Roll	X	The soundfield is "tilted" clockwise. For instance, a sound that starts at the front does not move, a sound that starts at the left moves upwards and a sound directly above moves right.

3.25.2.6 Ordering

This control determines the order in which the yaw, pitch and roll rotations are applied to make up a single rotation. The default is Y-P-R (i.e. yaw, pitch and then roll) but this can be switched to R-P-Y (i.e. roll, pitch and then yaw).

3.25.2.7 LPF Cutoff

This controls the cutoff frequency of a simple -6dB/octave low-pass filter applied to the delay output. This can be used to make the sound more distant as echoes repeat. Feedback can result in the filter being applied more than once to the same material.

3.25.2.8 Clear Delay

The "Clear Delay" button can be found at the bottom right of the plugin. While pressed, this button clears the delay line and silences everything but the dry signal.

4 Technical Notes

4.1 What is an O3A Stream?

A third order ambisonic (O3A) stream is made up of 16 individual channels of audio which together represent a 3D soundfield. Into this "audio scene" can be placed individual sound sources, reverberation and complex spatial textures.

These 16 channels can be quite confusing to understand conceptually. For instance, they do not relate to particular speaker directions, or to individual sound sources in the soundfield. However - you do *not* need to understand them to use them!

Each channel adds spatial detail to a sound scene. With just the first channel, you have a basically omnidirectional (mono) sound image. The second channel adds some basic detail left/right, the third up/down and the fourth front/back; these four channels make up first order ambisonics. The first and second channels together provide essentially the spatial detail available with the Mid/Side (M/S) stereo recording technique, which captures a sound image with left/right width. But with the further two channels the detail is available in all directions, not just left/right.

That covers the first 4 channels. The other 12 add further detail to make the image make sharper. If you are interested in what exactly is in these channels (which is not so easy to describe) you may want to read up on Higher Order Ambisonics (HOA). But you should *not* need to read up to *use* the techniques described here. If you want to make sense of what is going on spatially in an O3A stream, we find that it is normally best to use the [O3A Visualiser](#) plugin.

4.2 What processing can I apply to an O3A stream?

As well as processing designed specifically for O3A streams, it is possible to mix them together in the expected way. It is also possible to run conventional mono DSP algorithms on them directly by applying the algorithm to all 16 channels individually, subject to some rules. Failing to follow these rules is likely to shred the spatial imaging, so be careful! Specifically:

- If you process an O3A stream with a conventional mono DSP algorithm you *must apply the same processing to all 16 channels identically*.
- Only *linear* processing can be used (e.g. not distortion or compression). Also, be aware that time-variant processing can cause some issues.

If you are [working in Reaper](#), current versions do not make it particularly easy to set up this sort of processing, but it is possible if you are prepared to use some advanced features and do some rather tedious set-up. For instance, you can save an "FX Chain" in which eight stereo equalization plugins are "routed" correctly and in which controls are linked by "parameter modulation" so that one set of controls operates the others. Once the FX Chain is saved, you can load it into other projects.

4.3 Encoding

There are a number of ways in which the channels can be defined in HOA. To a large extent it does not matter which is chosen, as long as *everything uses the same convention*. This is critical and horrible things will happen to the spatial image if this is not the case. However, if you follow the convention, or convert explicitly where you need to, you will be able to pass audio around between different software packages.

These plugins use the "SN3D" ambisonic convention. As is usually the case, we order the channels using "ACN" numbering (see <http://ambisonics.ch/standards/channels>).

This encoding is used in the "AmbiX" file format and YouTube. It is supported directly by Rapture3D Universal.

4.4 Changes in Version 2.0

Please be aware that in version 2.0 of these plugins, we changed the ambisonic convention we use. We had been using the FuMa convention described below. Now, after consultation with our users, we are using SN3D/ACN which is future-proof and has been widely adopted.

4.5 How does O3A SN3D relate to FuMa and Classic Ambisonics?

The ambisonic format used (SN3D) is not directly compatible with classic 1970s four channel "WXYZ" B-Format, or the extension (FuMa) which was used by versions of these plugins prior to version 2.0.

However, conversion is straightforward. The free O3A Core pack includes plugins to do this: [O3A Decoder - FuMa](#) and [O3A Injector - FuMa](#).

4.6 Exporting First and Second Order

Third order ambisonics uses sixteen channels. To convert to first or second order, just take the first four or nine output channels and simply ignore the rest. This applies to both SN3D and FuMa.

Alternatively, the [O3A Decoder - O1A or O2A plugins](#) can be used to reduce the order of the material. This is a rare case where a plugin will work correctly for both SN3D and FuMa.

4.7 Why Third Order?

These plugins are based on Higher Order Ambisonic (HOA) techniques, which in principle can run at any "order" of spatial detail. We could have made all these plugins work at different orders; once you have got as far as making the maths and DSP work at third order (16 channels) it is not hard to extend to (for instance) fifth order (36 channels). We decided to use third order as it provides a good balance between spatial detail, current practice and CPU load. And sticking to just one type of B-Format avoids all sorts of "wiring" headaches!

That said, if you have a serious need for even higher order algorithms, please get in touch.